

IMPERIAL

# Neural Language Models

03/12/2025

Shamsuddeen Muhammad  
Google DeepMind Academic Fellow,  
Imperial College London  
<https://shmuhammadd.github.io/>

Idris Abdulmumin  
Postdoctoral Research Fellow,  
DSFSI, University of Pretoria  
<https://abumafrim.github.io/>

# Today

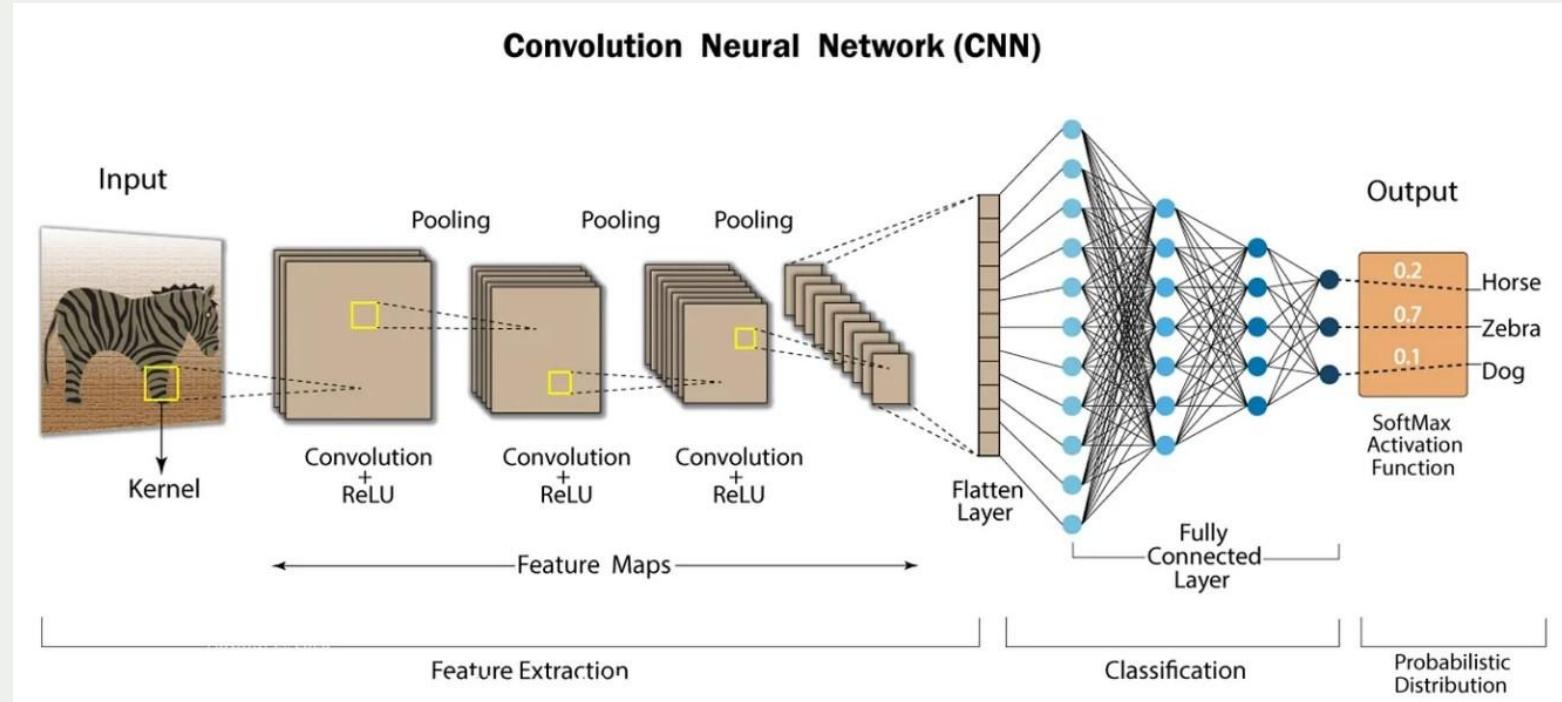
## Chapter 6 Neural Networks:

<https://web.stanford.edu/~jurafsky/slp3/6.pdf>

# Prerequisite

- Prerequisite
  - Neural Network
  - Loss function, backpropagation
  - CNN

# CNN Revisited



- A **CNN** is a type of neural network that's designed to work with *image data*.
- Robust Performance: Achieves state-of-the-art results in computer vision tasks.

# CNN Revisited

- They are made up of a series of layers:
  - The **convolutional** layers: responsible for extracting features from the images.
  - The **pooling** layers: responsible for reducing the size of the image after the convolutional layers have extracted the features.
  - **Activation** layers: Introduce non-linearity to the model (ReLU (Rectified Linear Unit), Sigmoid)

# CNN Revisited

In NLP, CNNs apply **1-dimensional convolutions** over sequences of word embeddings to automatically learn **local linguistic patterns** (e.g., n-grams such as “*not good*”, “*very happy*”)

Kim, Y. (2014). *Convolutional Neural Networks for Sentence Classification.*

# CNN Revisited

- Lack of Sequential Order Awareness
  - CNNs treat input as spatially ordered (like pixels in an image) but not temporally ordered.
  - In sequential tasks (e.g., text or time-series data), the order of input is critical.

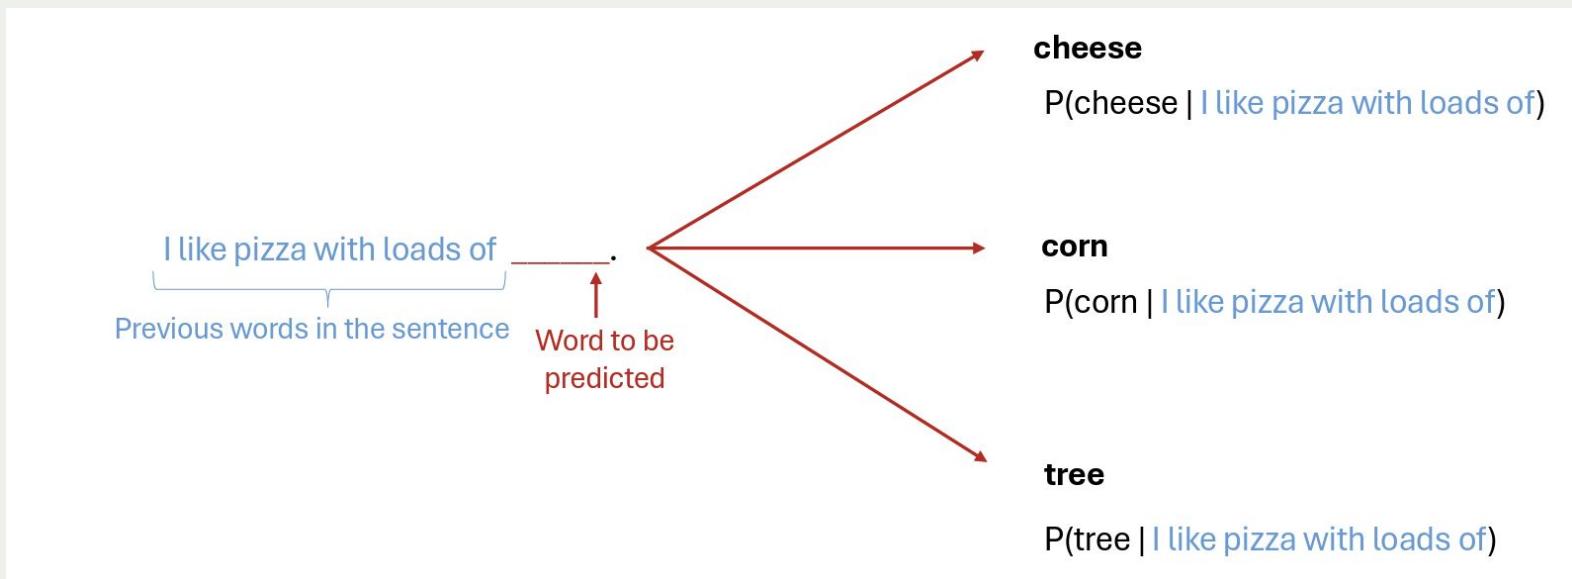
# CNN Revisited

CNNs require the input data to have a **fixed size** (e.g., images of a specific resolution).

This makes them unsuitable for sequences of varying lengths.

# Recall

A **language model (LM)** is a system that **predicts the next word** (or sequence of words) in a sentence.



# Recall: Language Modeling (LM)

**Goal:** compute the probability of a sentence or sequence of words  $W$ :

$$P(W) = P(w_1, w_2, w_3, w_4, w_5 \dots w_n)$$

Related task: probability of an upcoming word:

$$P(w_5 | w_1, w_2, w_3, w_4) \text{ or } P(w_n | w_1, w_2 \dots w_{n-1})$$

A language model computes either of these:

$$P(W) \text{ or } P(w_n | w_1, w_2 \dots w_{n-1})$$

# Categories for Language Models

- Statistical Language Models
- Neural Language Models (NNLMs)
- Transformer-Based Language Models
- Large Language Models (LLMs)
- Multilingual & Low-Resource LMs

# Limitation of N-Gram Language Model

- N-gram models struggle to capture **long-range dependencies** in natural language.
- They only look at the previous  $n$  words, so important relationships across the sentence are lost.
- They lack global coherence and fail on complex syntax.

# Traditional Models

The following classical approach treats text as unordered.

- Rule-Based & Pattern Matching Systems
- Classical Machine Learning : Logistic Regression, Naive Bayes, SVM (Support Vector Machines), Decision Trees / Random Forests
- TF-IDF + Linear Classifiers (Logistic Regression, SVM)
- Topic Models such as LDA

# Limitations of Classical Models

- Ignore word order (bag-of-words assumption)
- Cannot capture long-range dependencies
- Fail on tasks needing context
- Limited semantic understanding

## Simple Example for Students

Sentence A: “*Football makes people happy.*”

Sentence B: “*People make football happy.*”

Both contain the same words - BoW models treat them as identical, even though the meanings differ.

# Sequence Modelling

Sequence modelling refers to building models that can ***represent, process, and predict ordered data, where the order of elements matters.***

- Natural language is inherently sequential: meaning depends on previous and next words.
- Many NLP tasks cannot be solved by treating text as unordered bags of words.

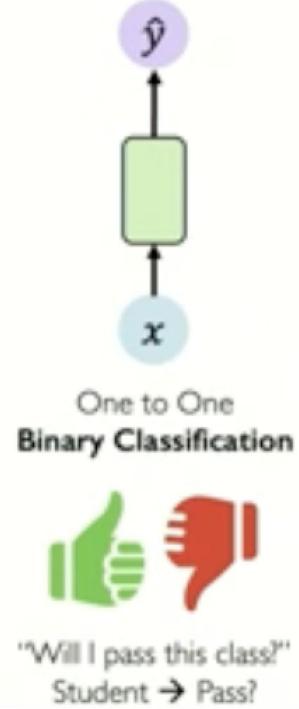


# Classical sequence architectures

Classical sequence architectures:

- Recurrent Neural Networks (RNNs)
- Long Short-Term Memory Networks (LSTMs)
- Gated Recurrent Units (GRUs)
- Encoder–Decoder architectures (precursor to Transformers)

# Sequence Modelling Application



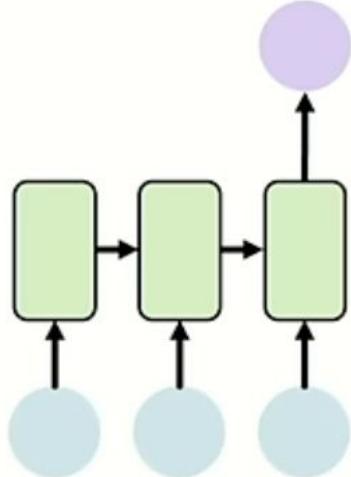
# Sequence Modelling Application



One to One  
**Binary Classification**



"Will I pass this class?"  
Student → Pass?



Many to One  
**Sentiment Classification**



Ivar Hagendoorn  
@IvarHagendoorn

Follow

The @MIT Introduction to #DeepLearning is definitely one of the best courses of its kind currently available online [introtodeeplearning.com](http://introtodeeplearning.com)

12:45 PM - 12 Feb 2018



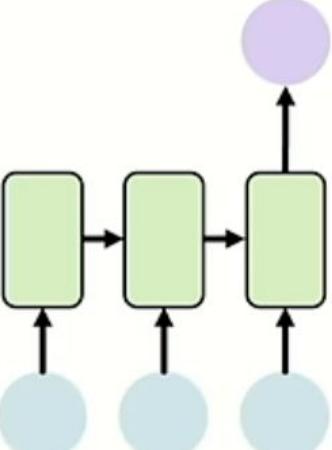
# Sequence Modelling Application



One to One  
**Binary Classification**



"Will I pass this class?"  
Student → Pass?



Many to One  
**Sentiment Classification**

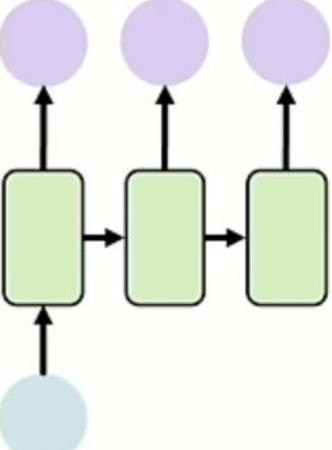


Ivar Hagendoorn  
@IvarHagendoorn



The @MIT Introduction to #DeepLearning is definitely one of the best courses of its kind currently available online [introtodeeplearning.com](http://introtodeeplearning.com)

12:45 PM - 12 Feb 2018

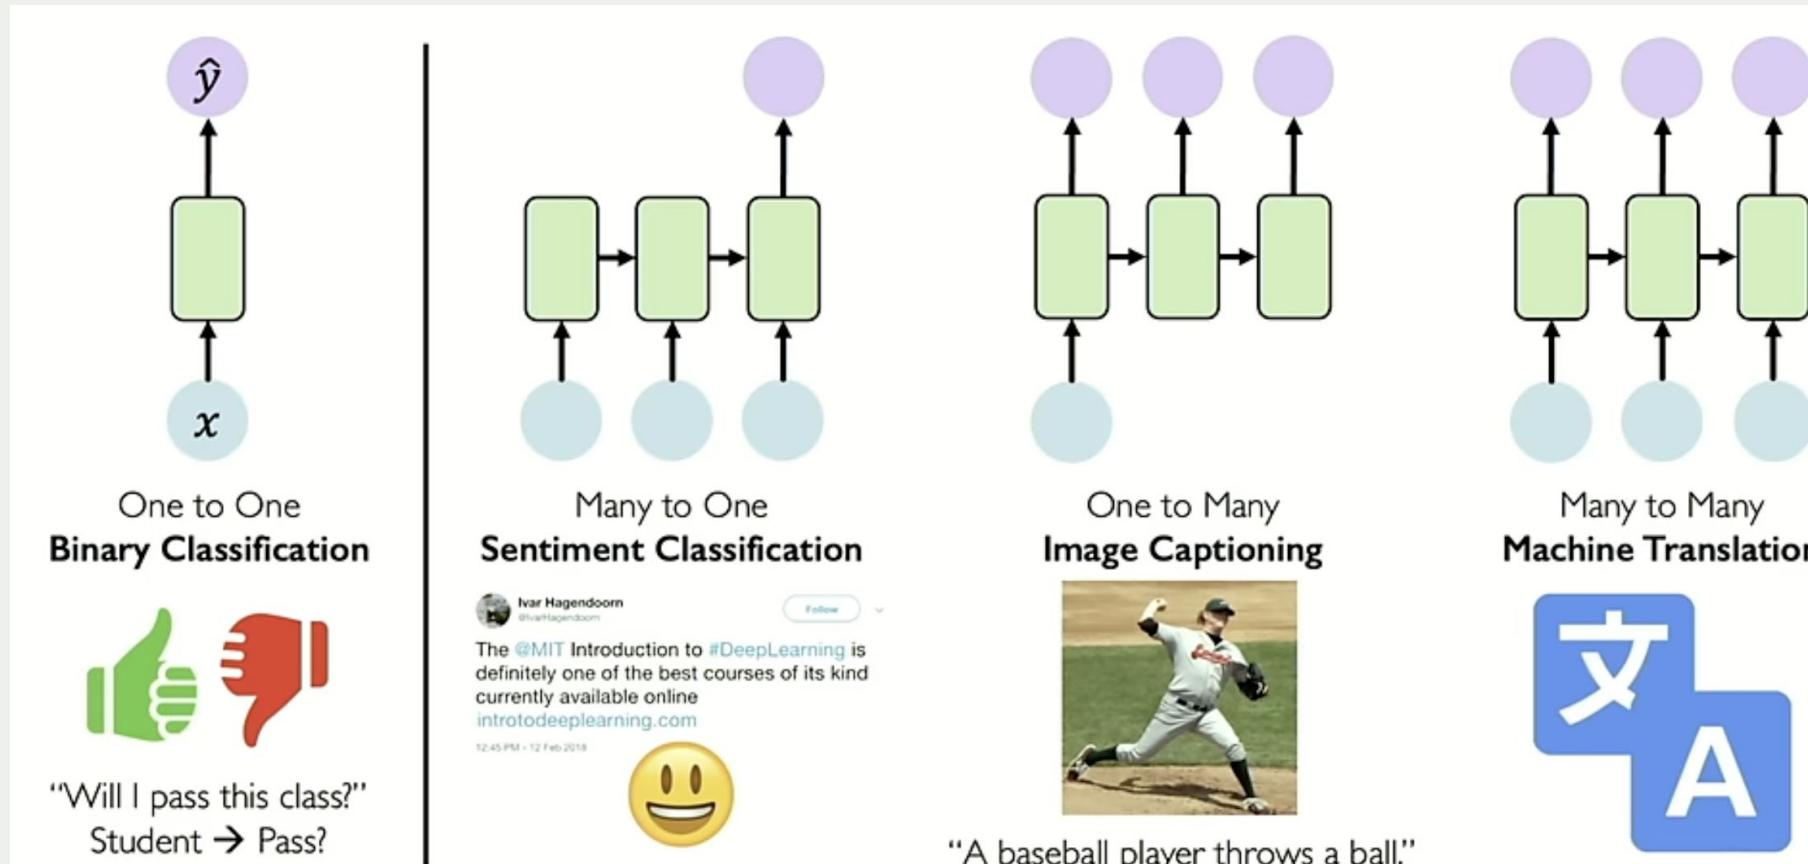


One to Many  
**Image Captioning**



"A baseball player throws a ball."

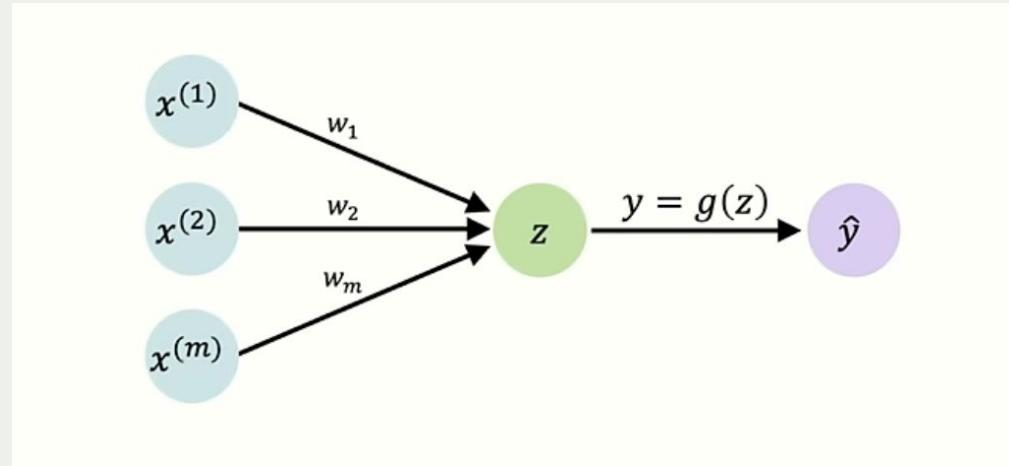
# Sequence Modelling Application



# Neurons with Recurrence

# Perceptron Revisited

- The **Perceptron** is one of the earliest and simplest learning algorithms for binary classification.
- It learns a **linear decision boundary** by repeatedly adjusting its weight vector based on mistakes.

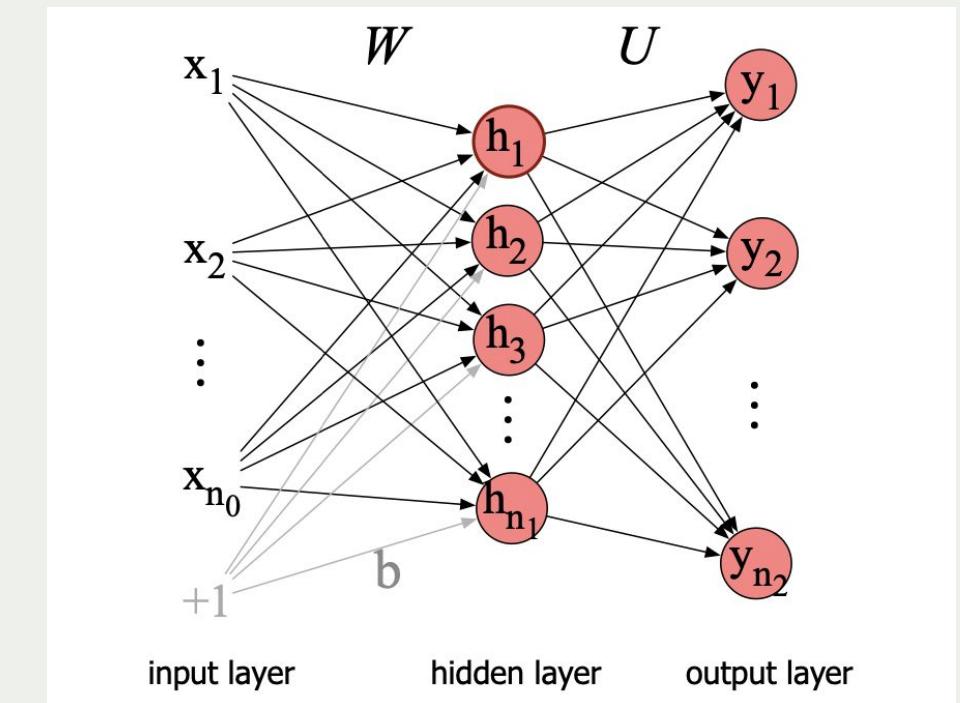


# Feed-Forward Network

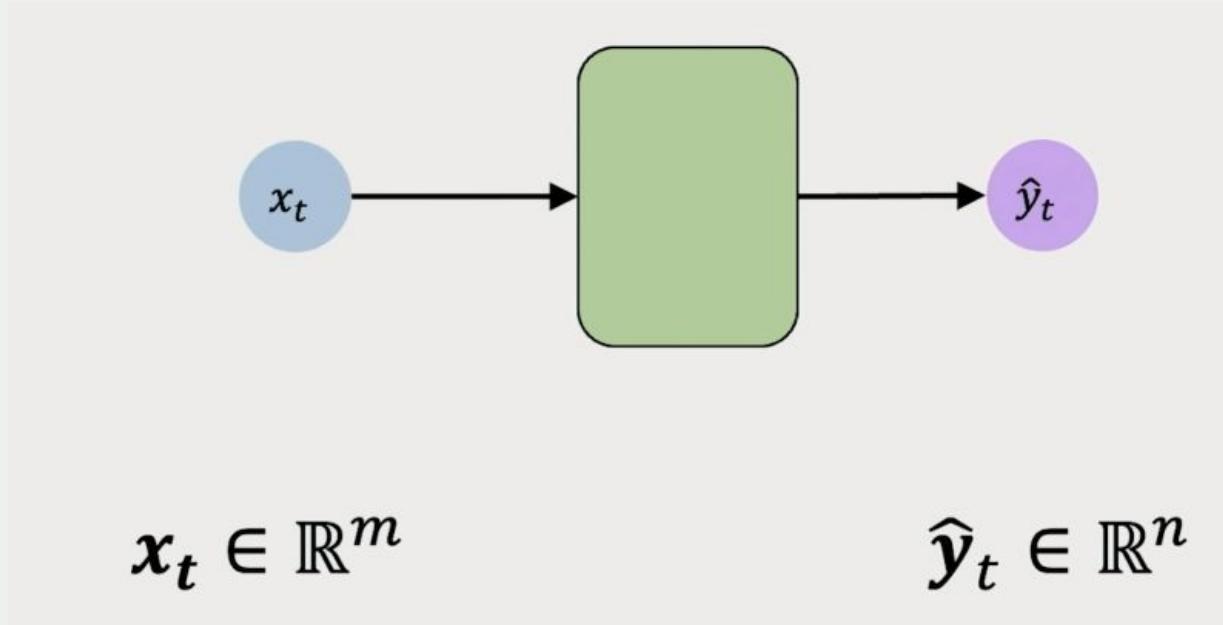
A **Feed-Forward Neural Network (FNN)**, also called a *multi-layer perceptron (MLP)*, is a basic neural architecture where information flows in one direction

**layers of perceptrons** → can model nonlinear decision boundaries and richer semantic relationships.

No notion of time step or sequence

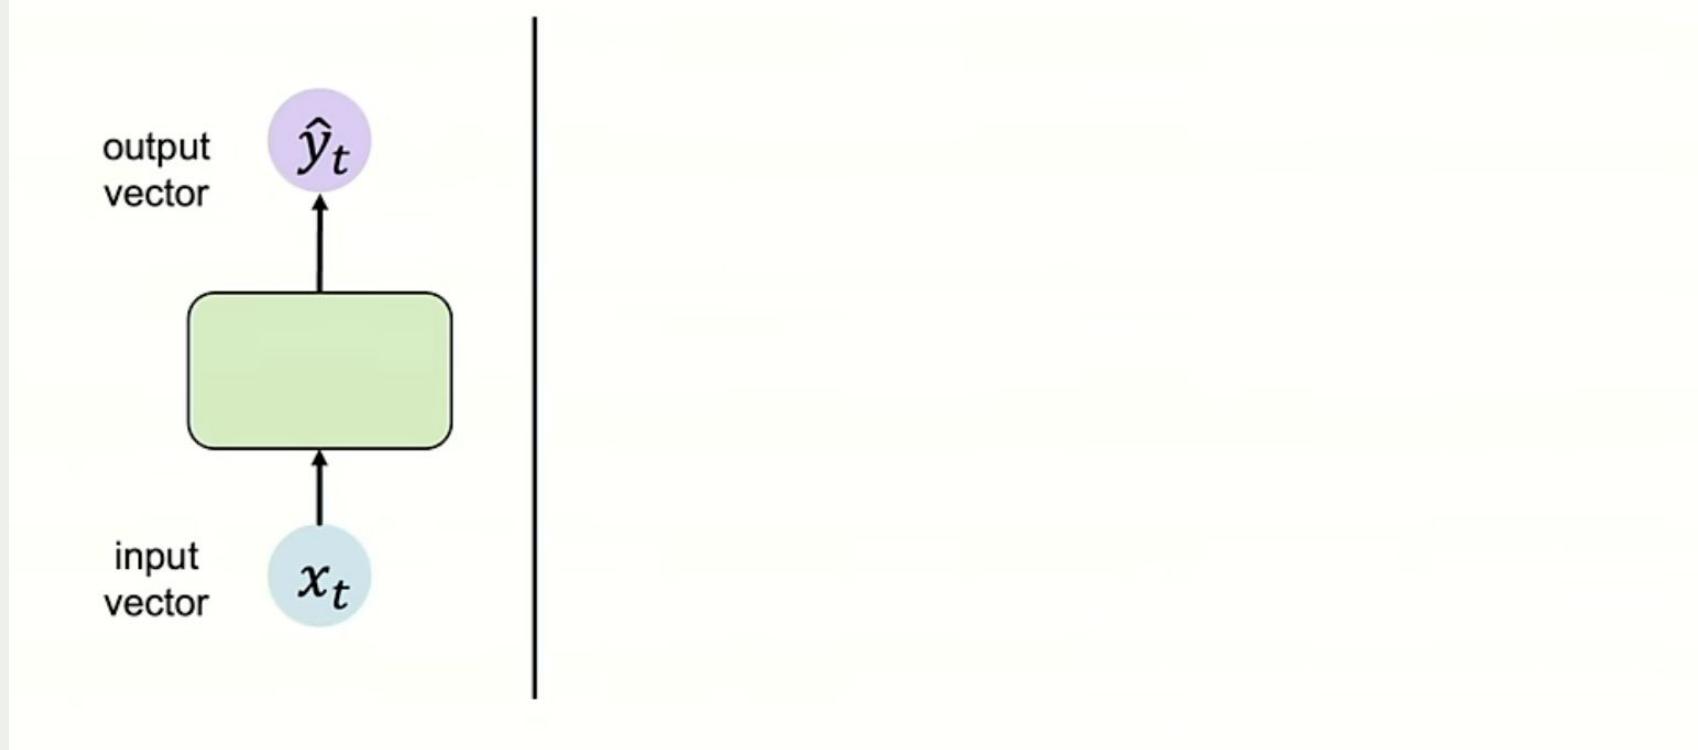


# Feed-Forward Network Revisited

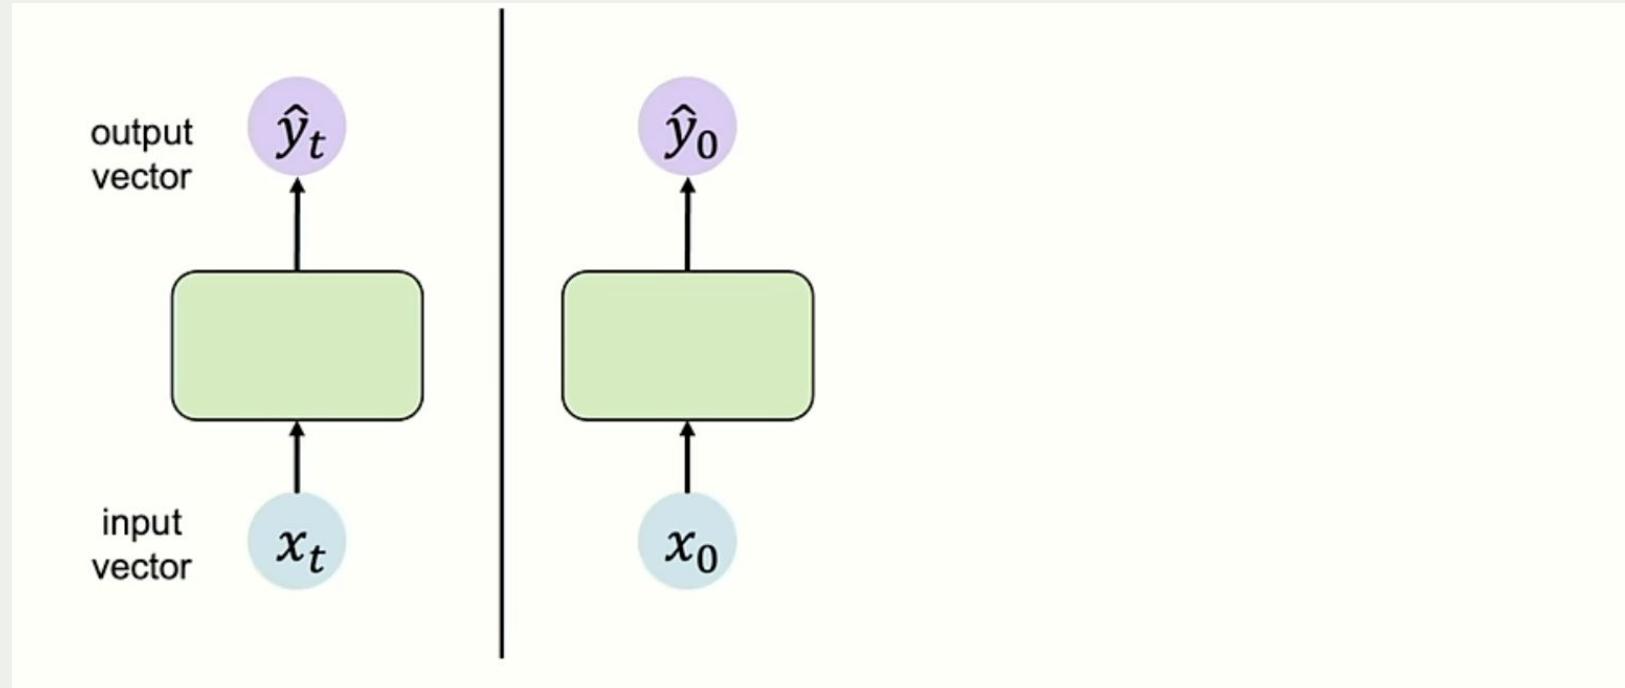


- Still no notion don't handle sequence data
- How can we handle sequence data?

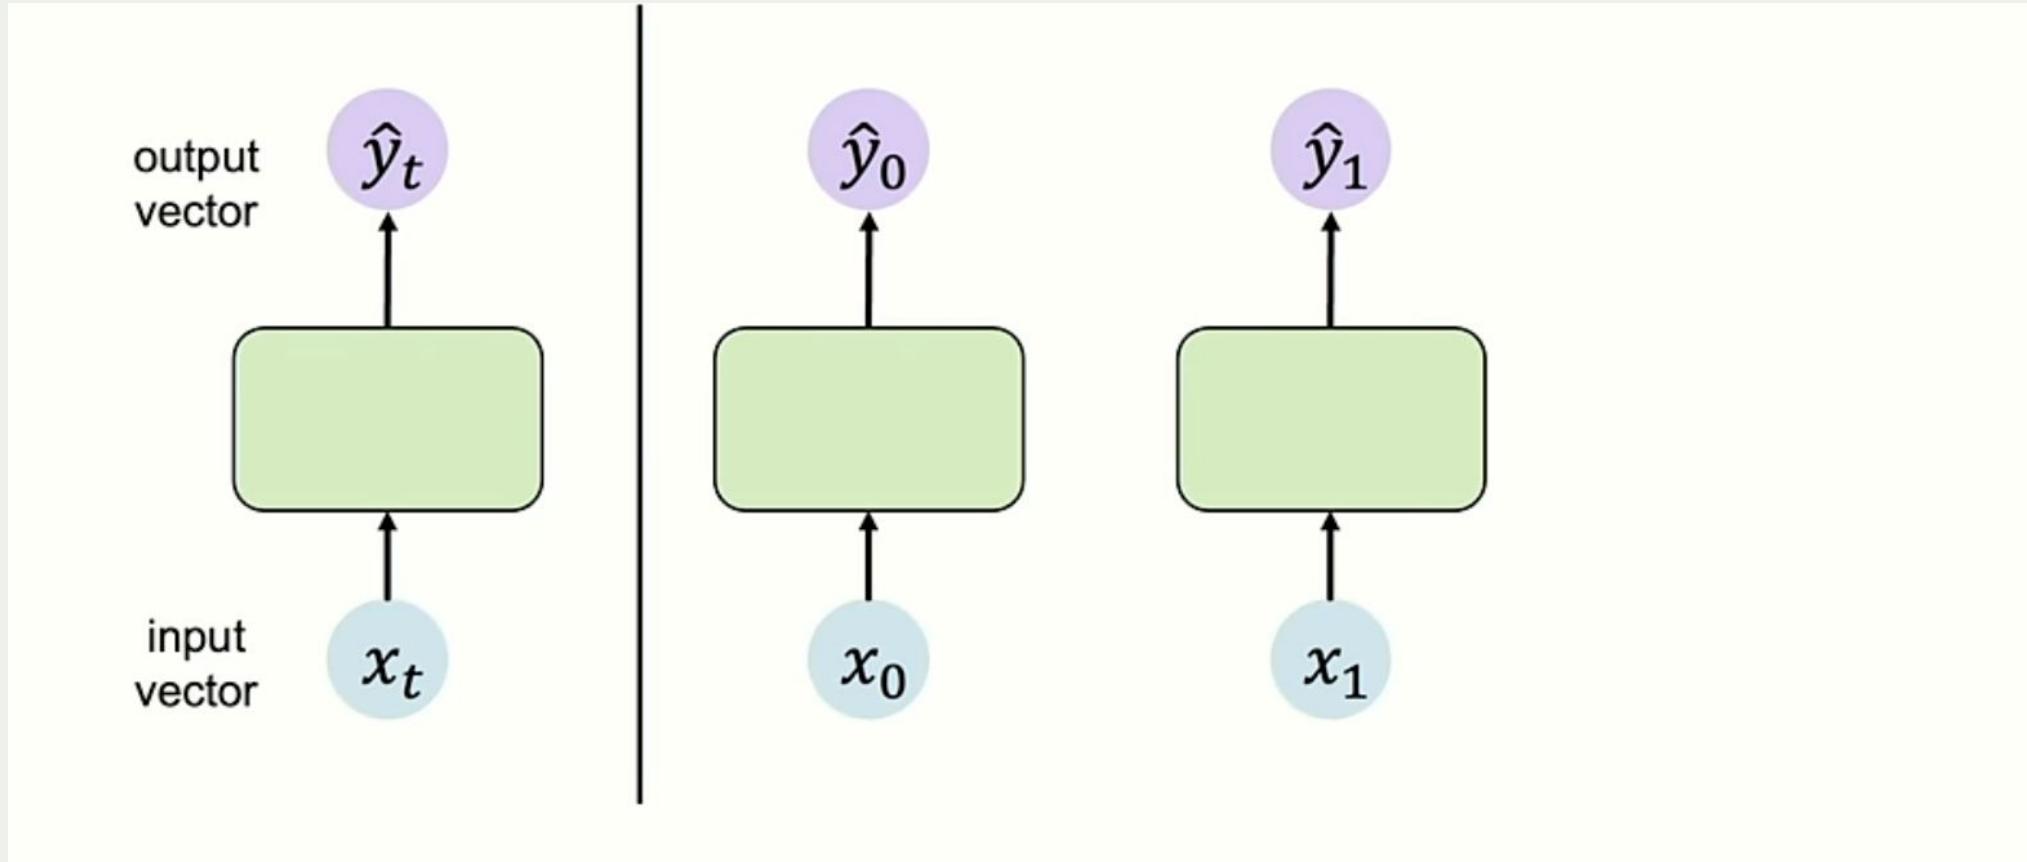
# Handling Individual Time-Step



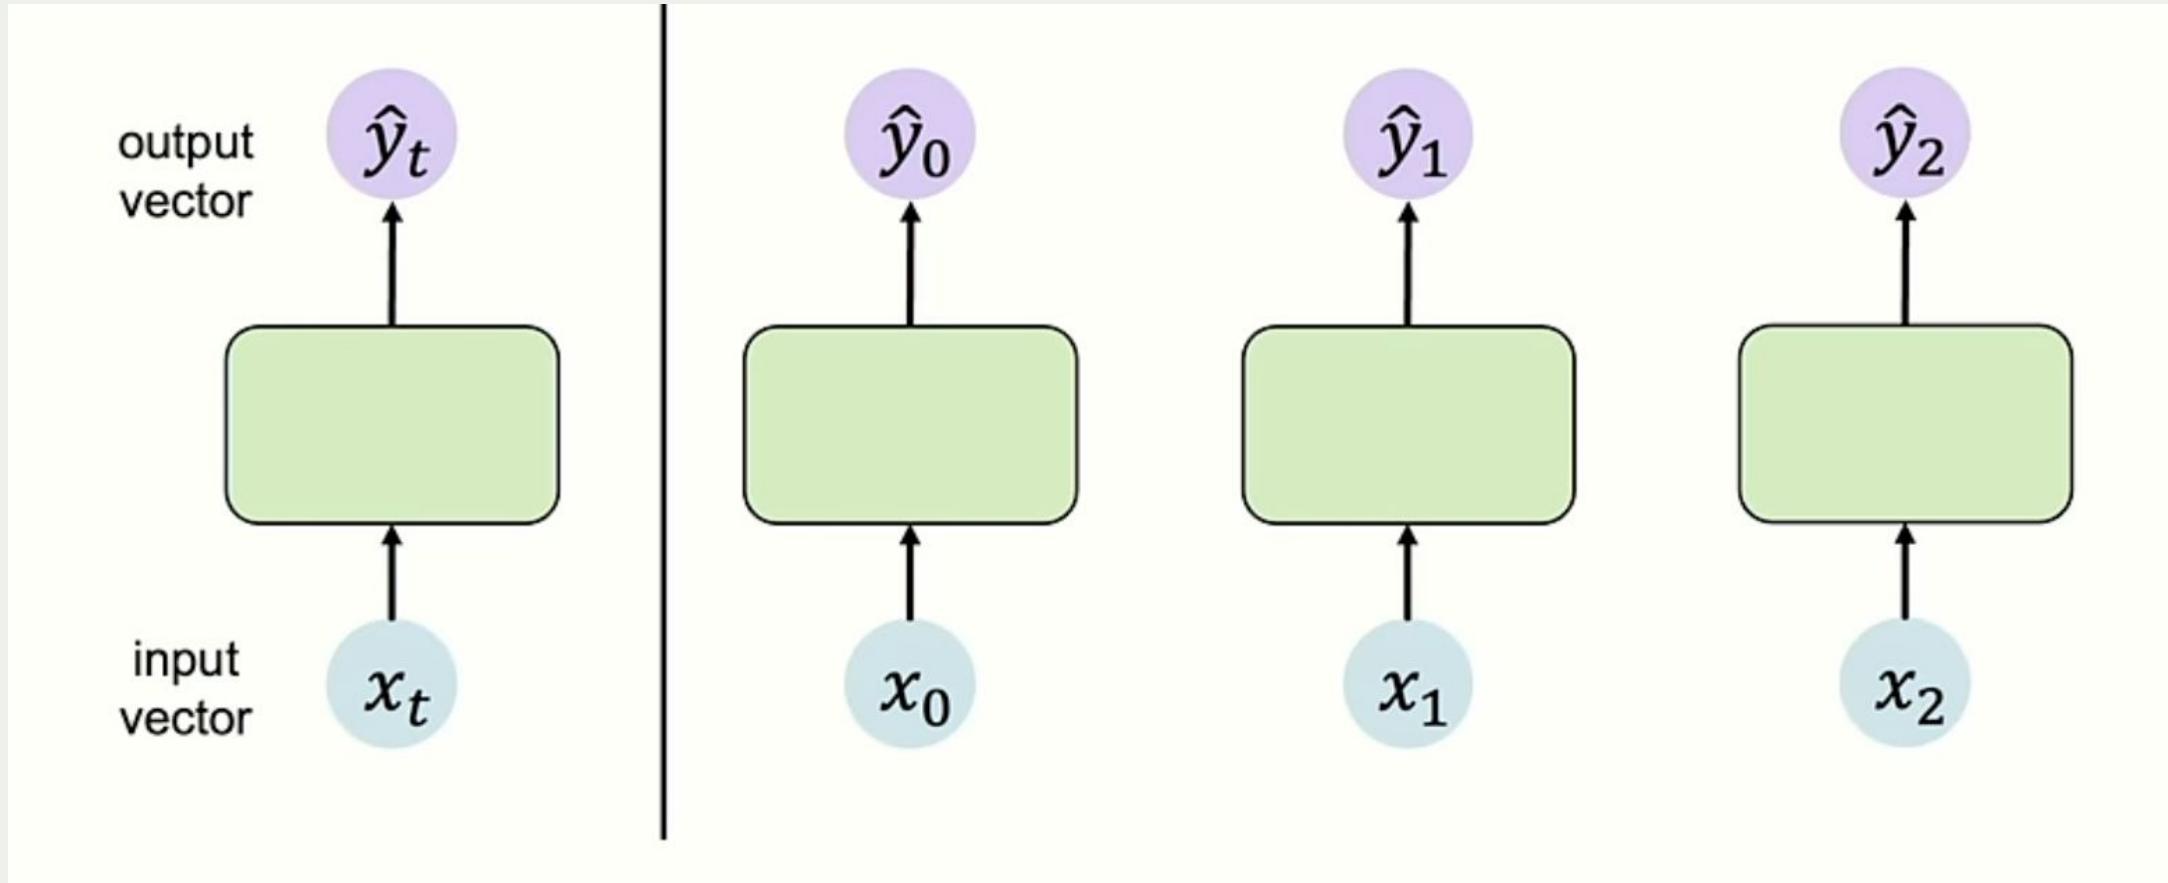
# Handling Individual Time-Step



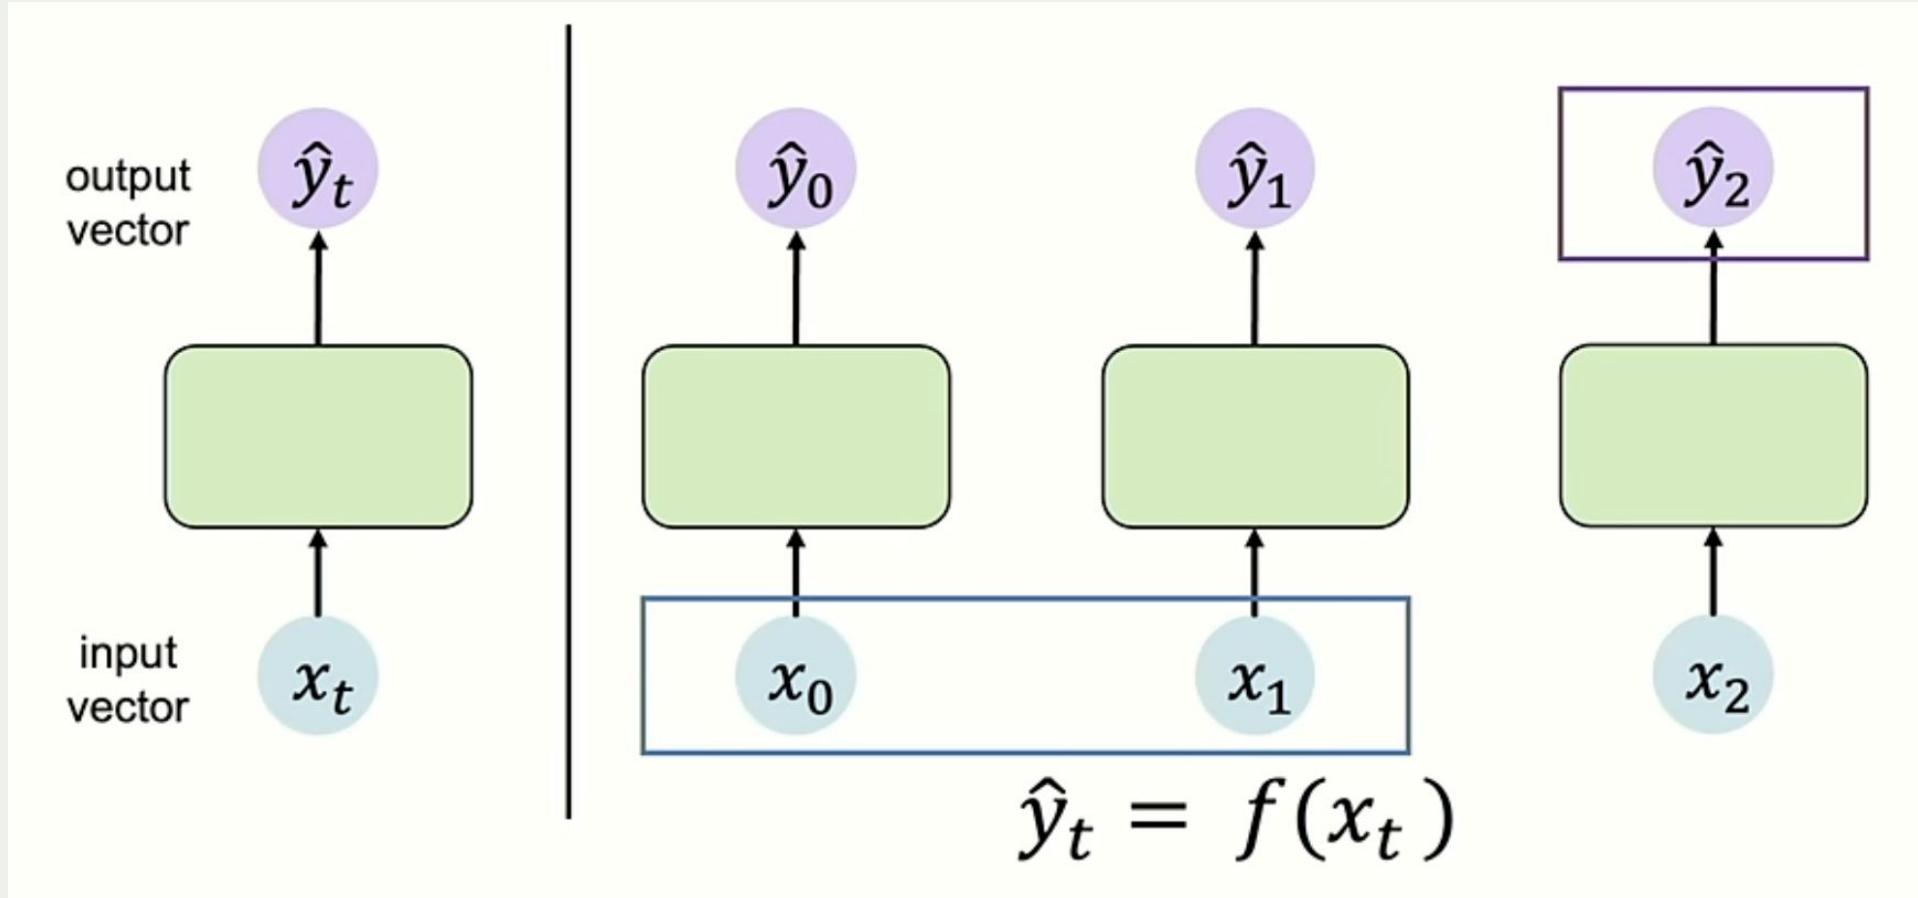
# Handling Individual Time-Step



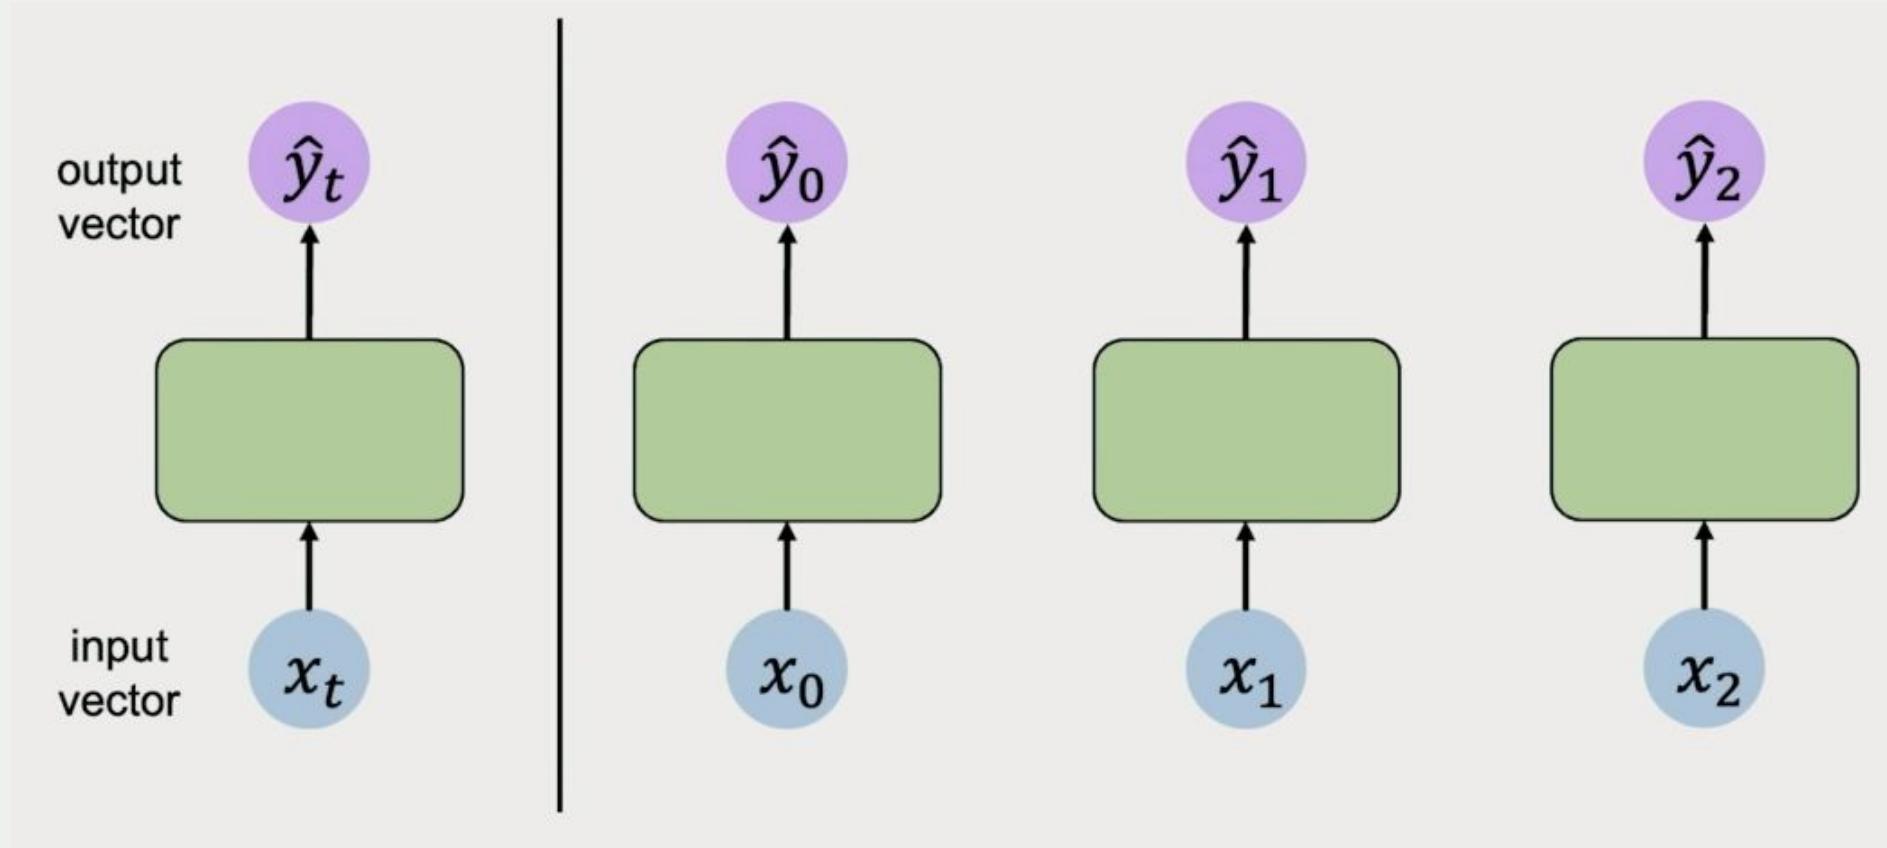
# Handling Individual Time-Step



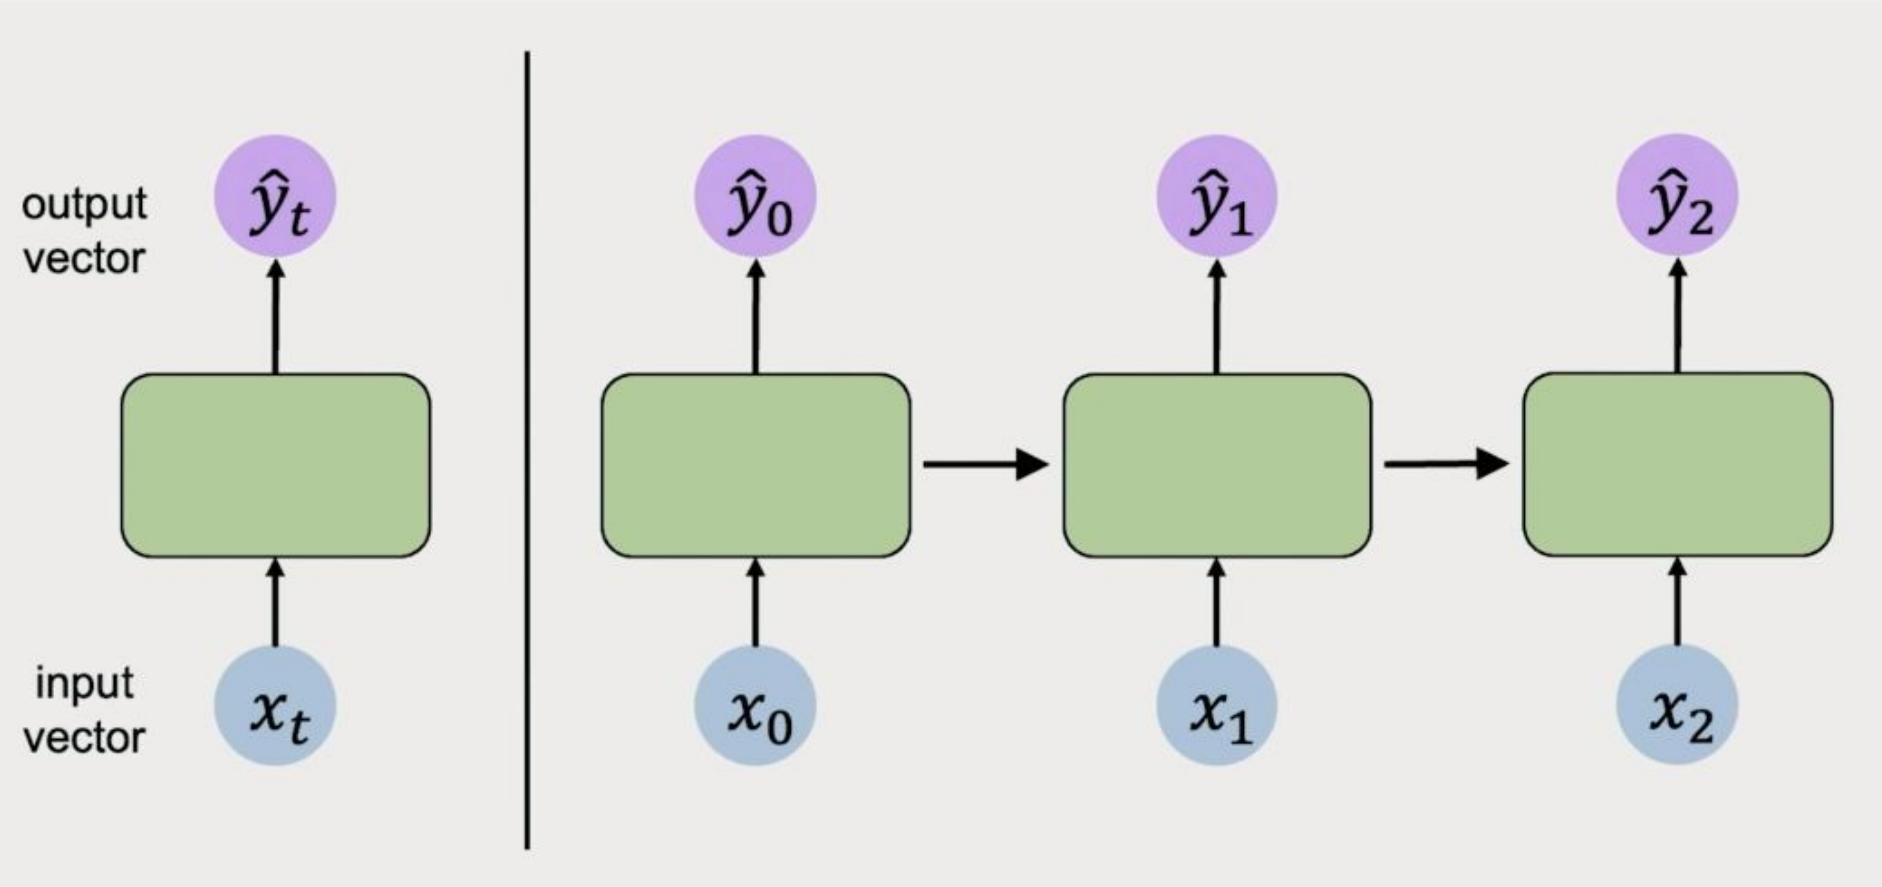
# Handling Individual Time-Step



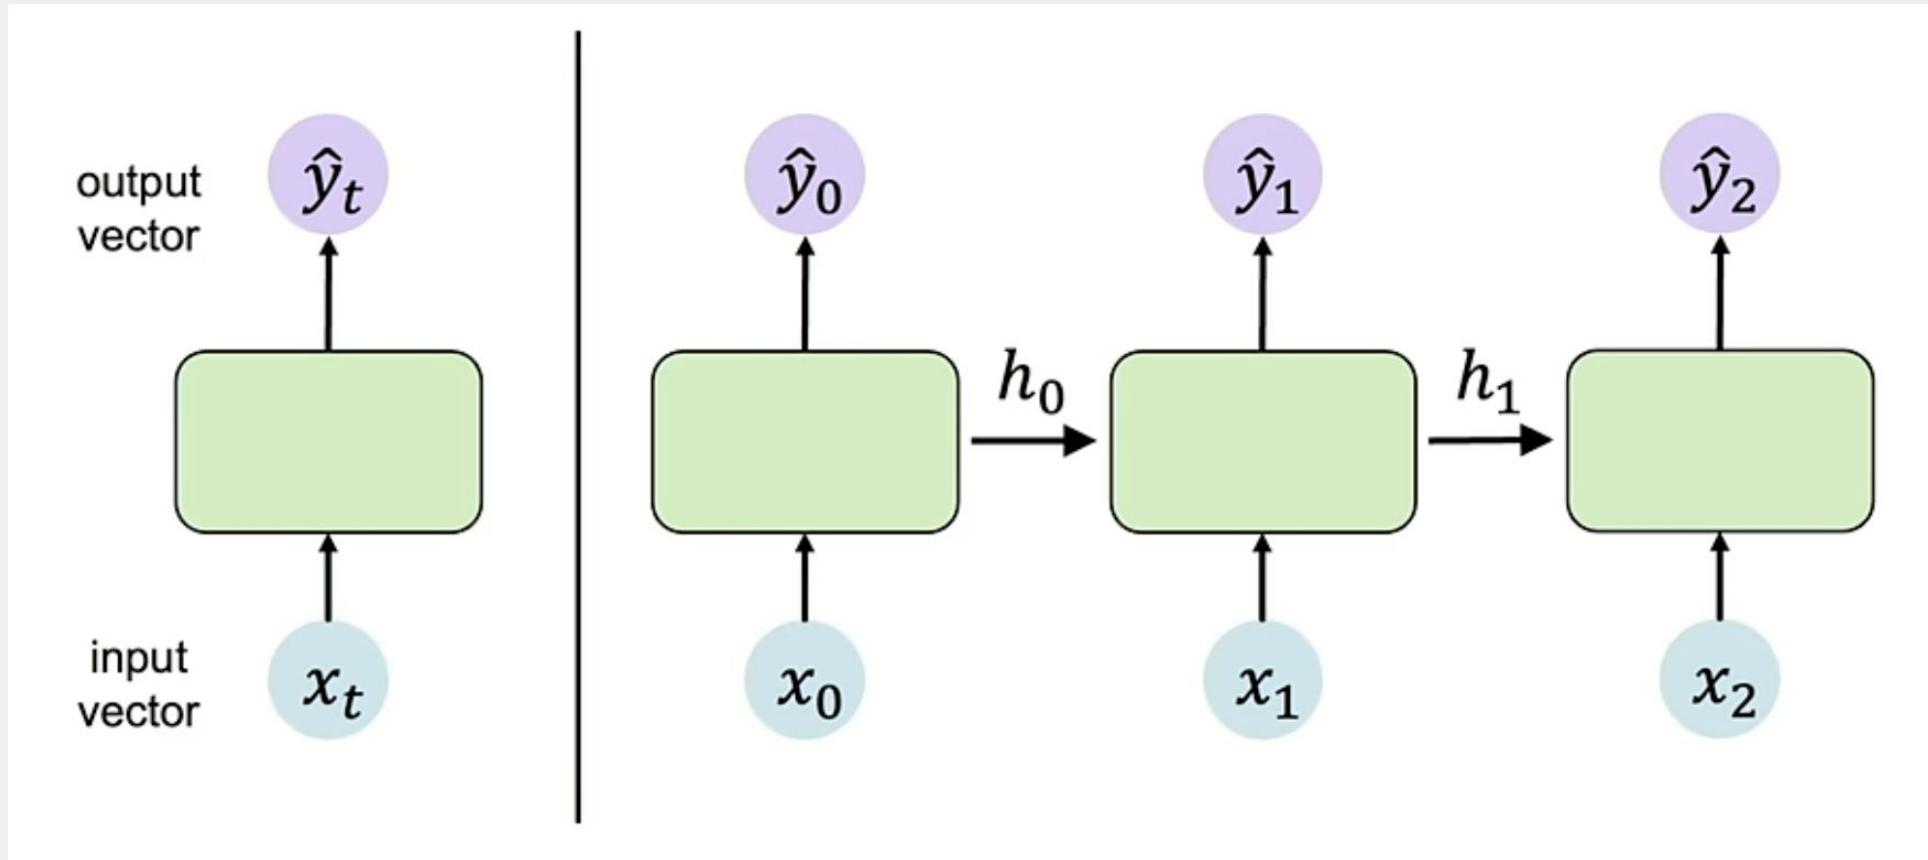
# Neurons with Recurrent



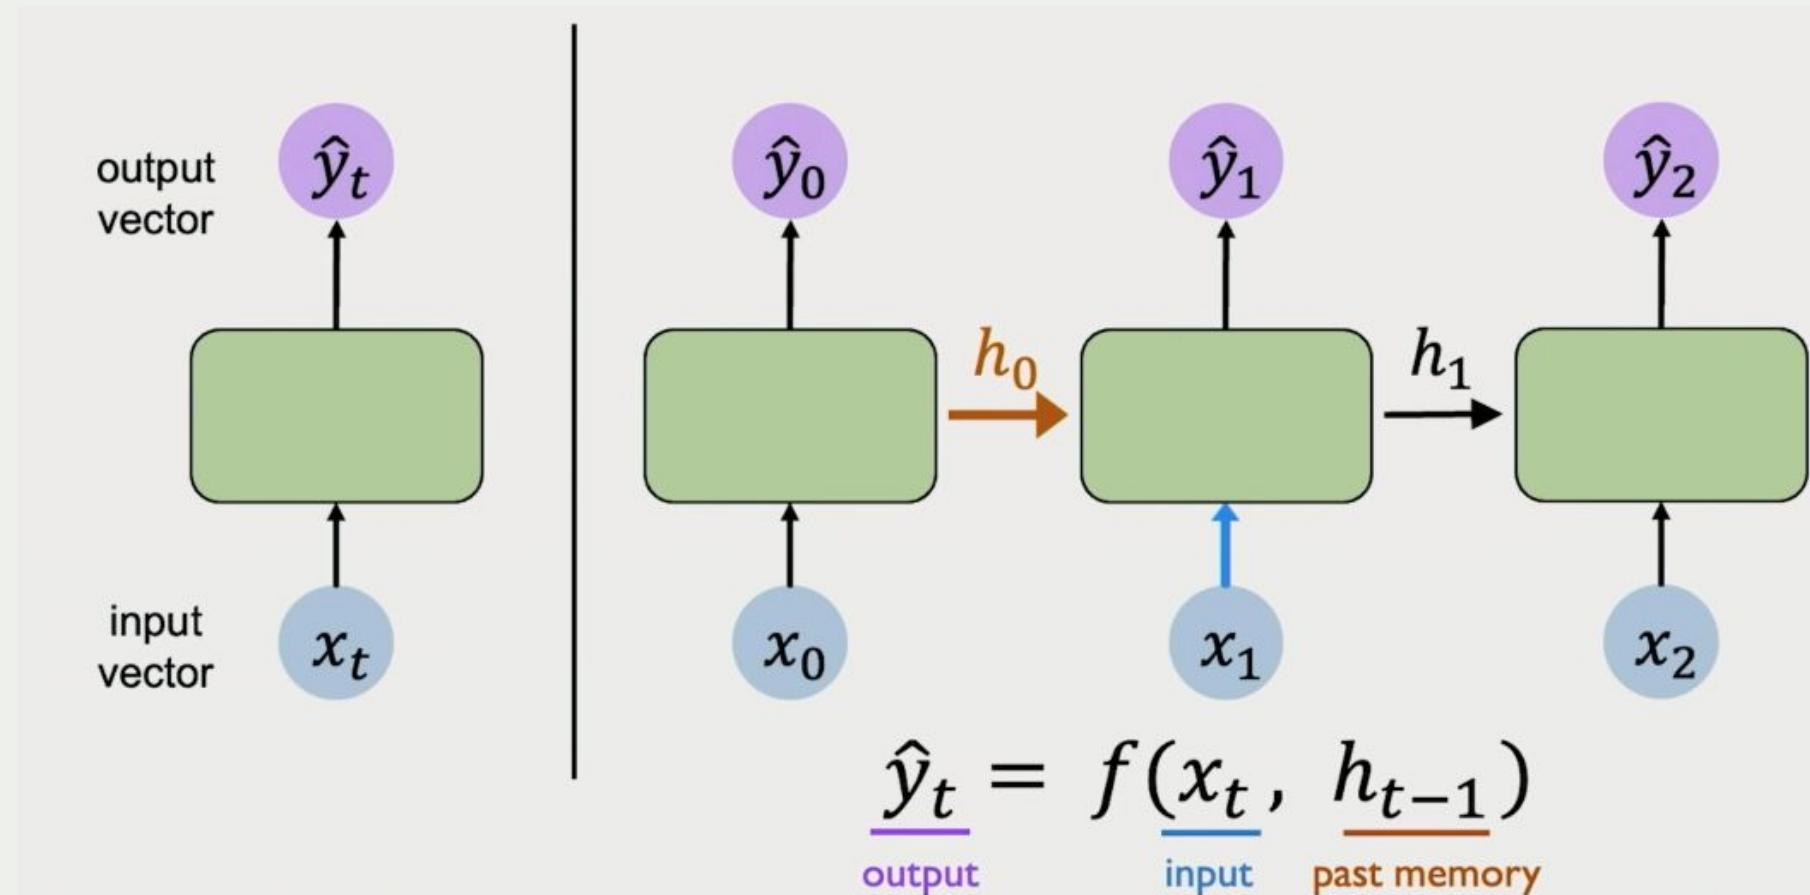
# Neurons with Recurrent



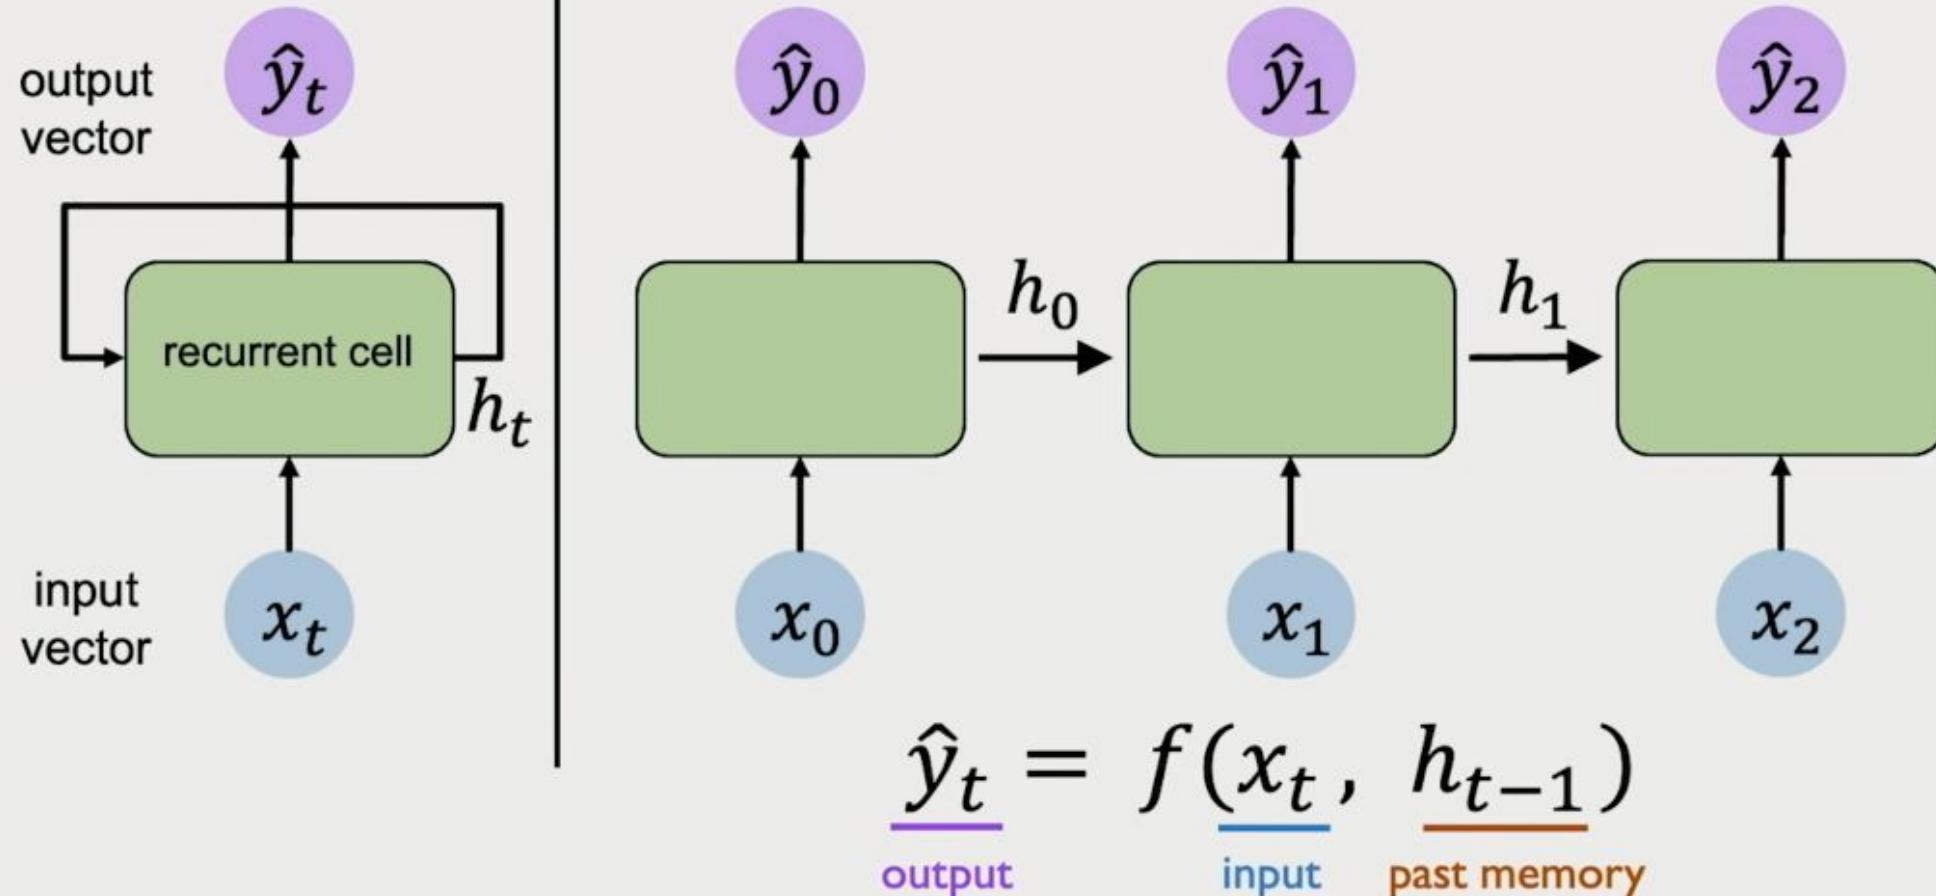
# Neurons with Recurrent



# Neurons with Recurrent



# Neurons with Recurrent



# Recurrent

The term "recurrent" generally means occurring repeatedly or in cycles, and in the context of Recurrent Neural Networks (RNNs), it specifically refers to the mechanism of repeatedly applying the same computational process to sequential data.

The recurrence enables the network to "remember" previous states over time, giving RNNs the ability to capture patterns and dependencies across sequential data.

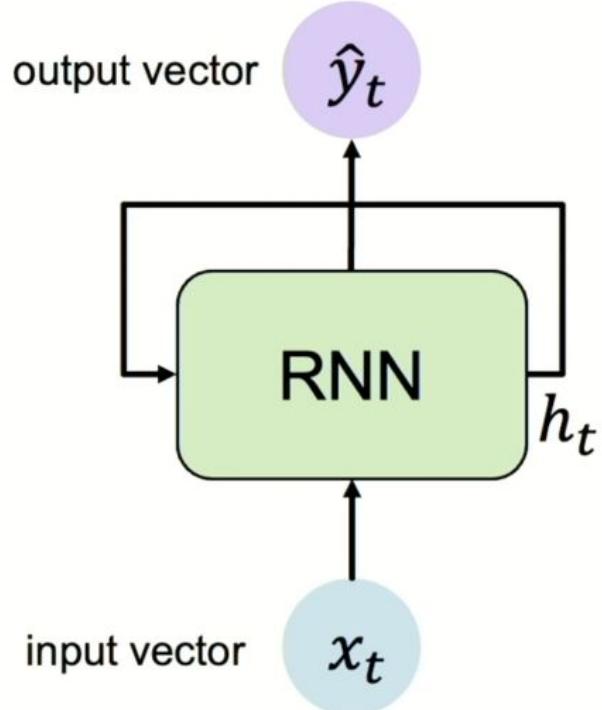
# Recurrent

The recurrence is foundational concept sequence modelling called Recurrent Neural Networks (RNN)

- **Text sequences** (e.g., sentences in natural language processing),
- **Time-series data** (e.g., stock prices, sensor readings),
- **Speech data** (e.g., audio signals),
- **Video frames** (e.g., in action recognition).

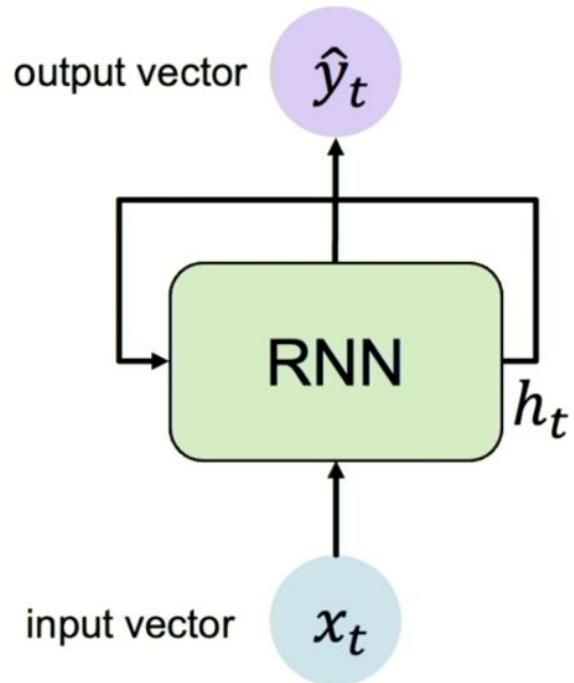
# Recurrent Neural Networks (RNNs)

# Recurrent Neural Network



RNNs have a **state**,  $h_t$ , that is updated **at each time step** as a sequence is processed

# Recurrent Neural Network



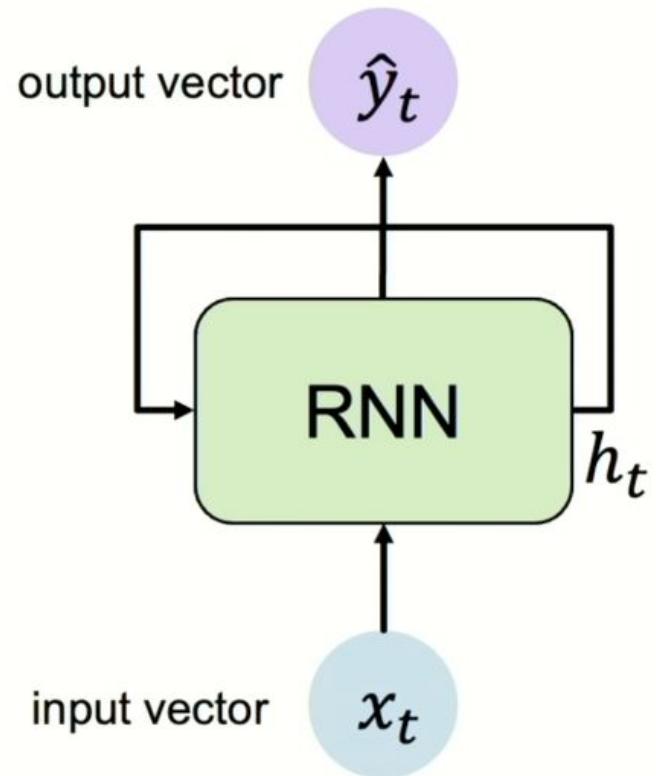
Apply a **recurrence relation** at every time step to process a sequence:

$$h_t = f_W(x_t, h_{t-1})$$

cell state

RNNs have a **state**,  $h_t$ , that is updated **at each time step** as a sequence is processed

# Recurrent Neural Network



Apply a **recurrence relation** at every time step to process a sequence:

$$h_t = f_W(x_t, h_{t-1})$$

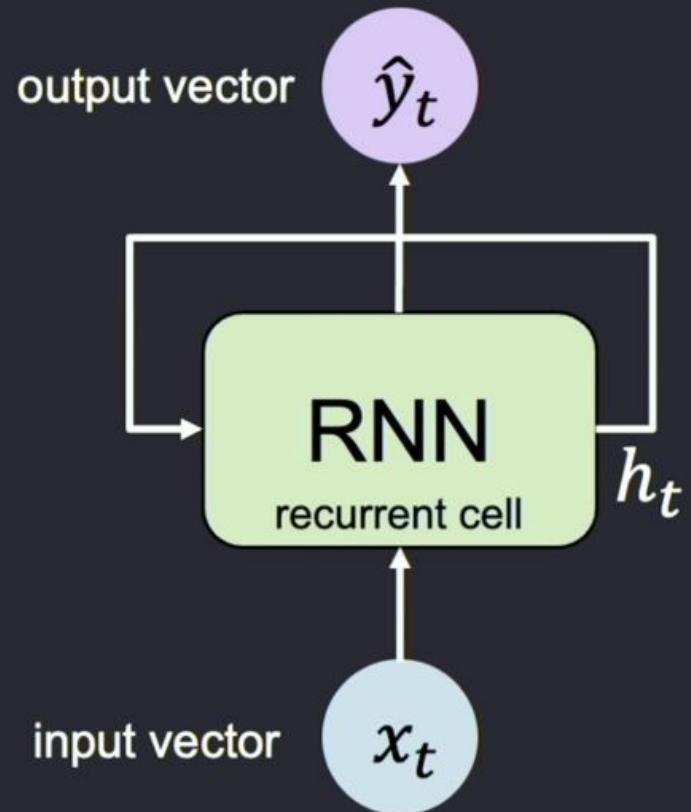
cell state      function with weights  $W$       input      old state

Note: the same function and set of parameters are used at every time step

RNNs have a **state**,  $h_t$ , that is updated **at each time step** as a sequence is processed

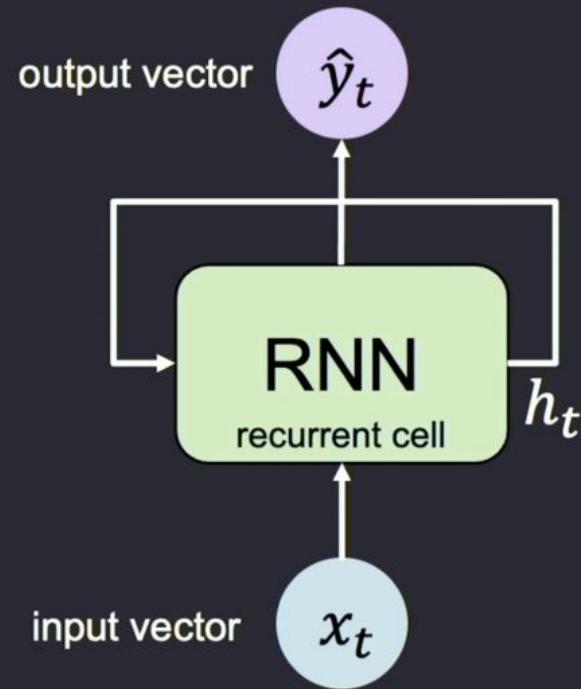
# RNN Intuition: Pseudocode

```
my_rnn = RNN()  
hidden_state = [0, 0, 0, 0]  
  
sentence = ["I", "love", "recurrent", "neural"]  
  
for word in sentence:  
    prediction, hidden_state = my_rnn(word, hidden_state)  
  
    next_word_prediction = prediction  
    # >>> "networks!"
```



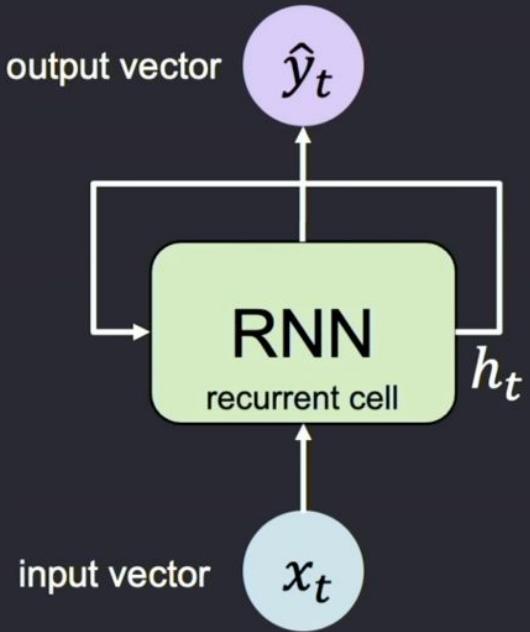
# RNN Intuition: Pseudocode

```
my_rnn = RNN()  
hidden_state = [0, 0, 0, 0]  
  
sentence = ["I", "love", "recurrent", "neural"]  
  
for word in sentence:  
    prediction, hidden_state = my_rnn(word, hidden_state)  
  
next_word_prediction = prediction  
# >>> "networks!"
```

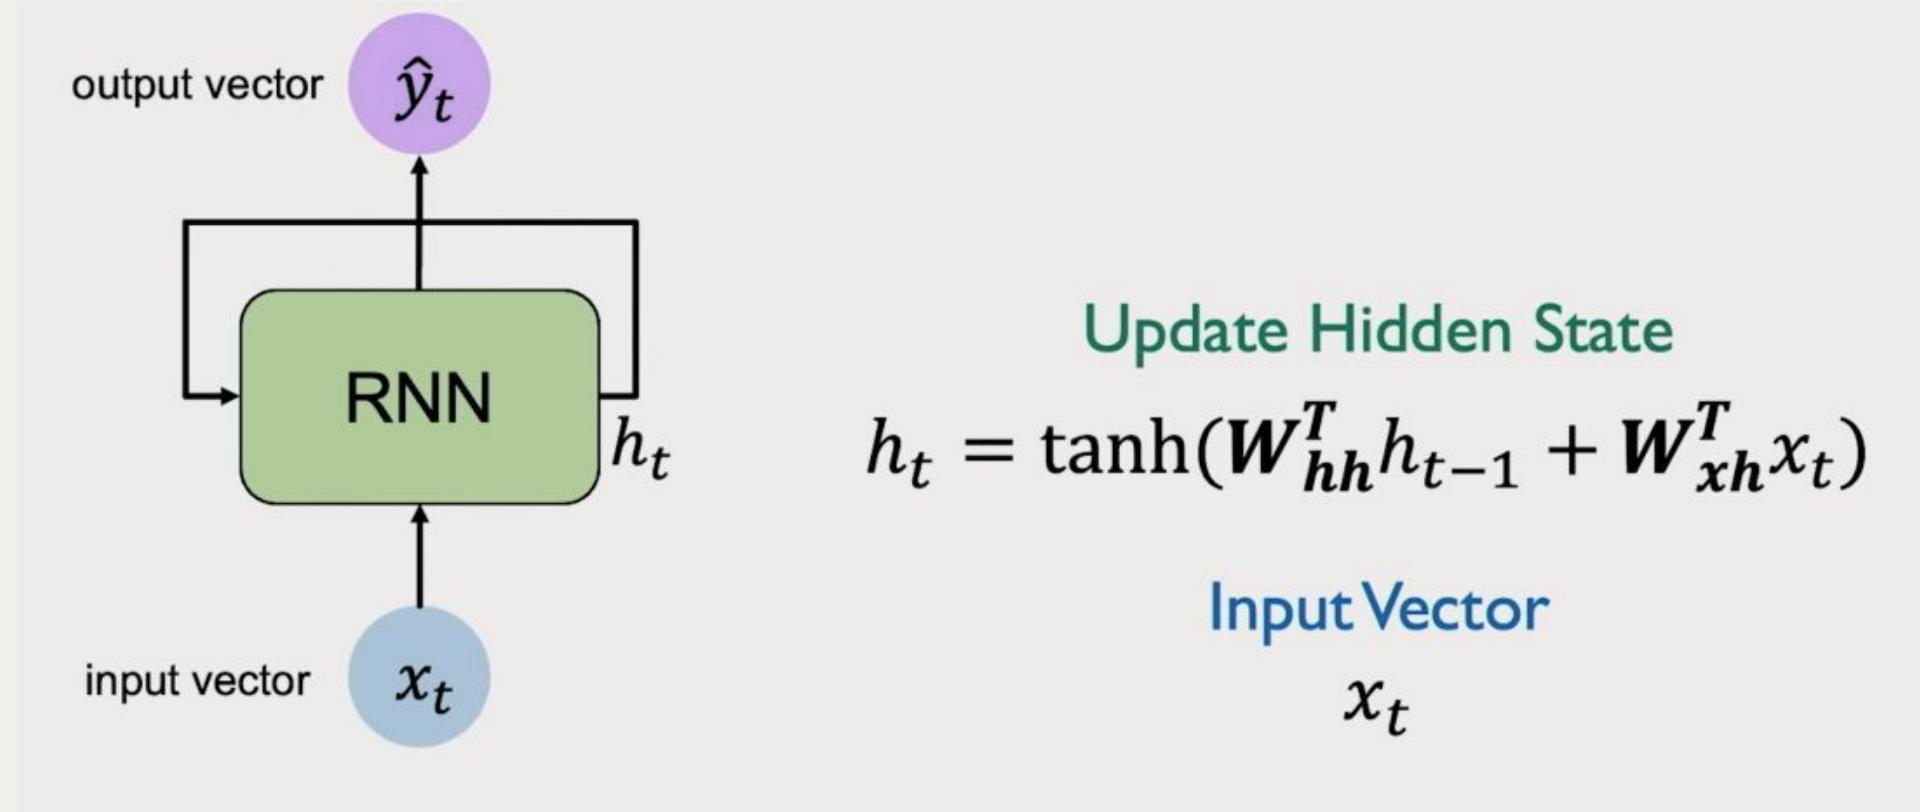


# RNN Intuition: Pseudocode

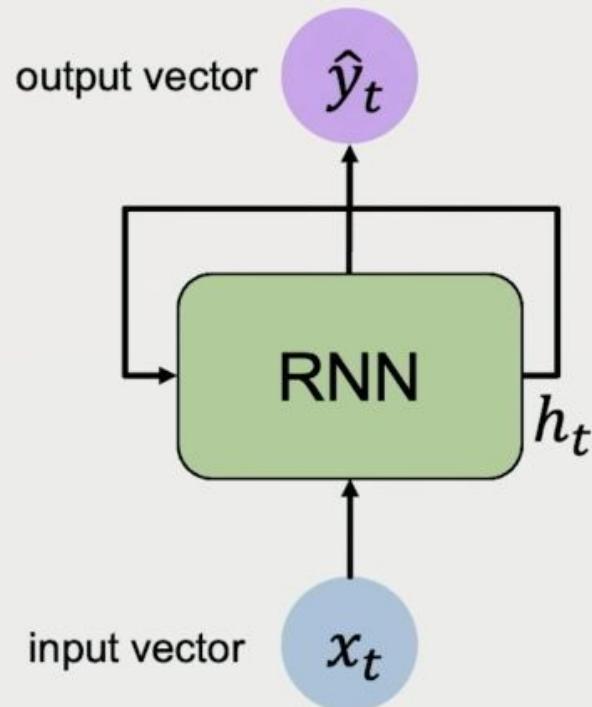
```
my_rnn = RNN()  
hidden_state = [0, 0, 0, 0]  
  
sentence = ["I", "love", "recurrent", "neural"]  
  
for word in sentence:  
    prediction, hidden_state = my_rnn(word, hidden_state)  
  
    next_word_prediction = prediction  
    # >>> "networks!"
```



# RNN State Update and Output



# RNN: Computational Graph Across Time



Output Vector

$$\hat{y}_t = \mathbf{W}_{hy}^T h_t$$

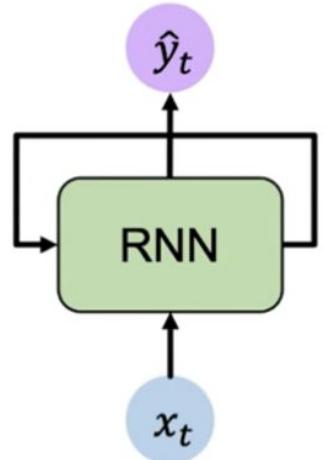
Update Hidden State

$$h_t = \tanh(\mathbf{W}_{hh}^T h_{t-1} + \mathbf{W}_{xh}^T x_t)$$

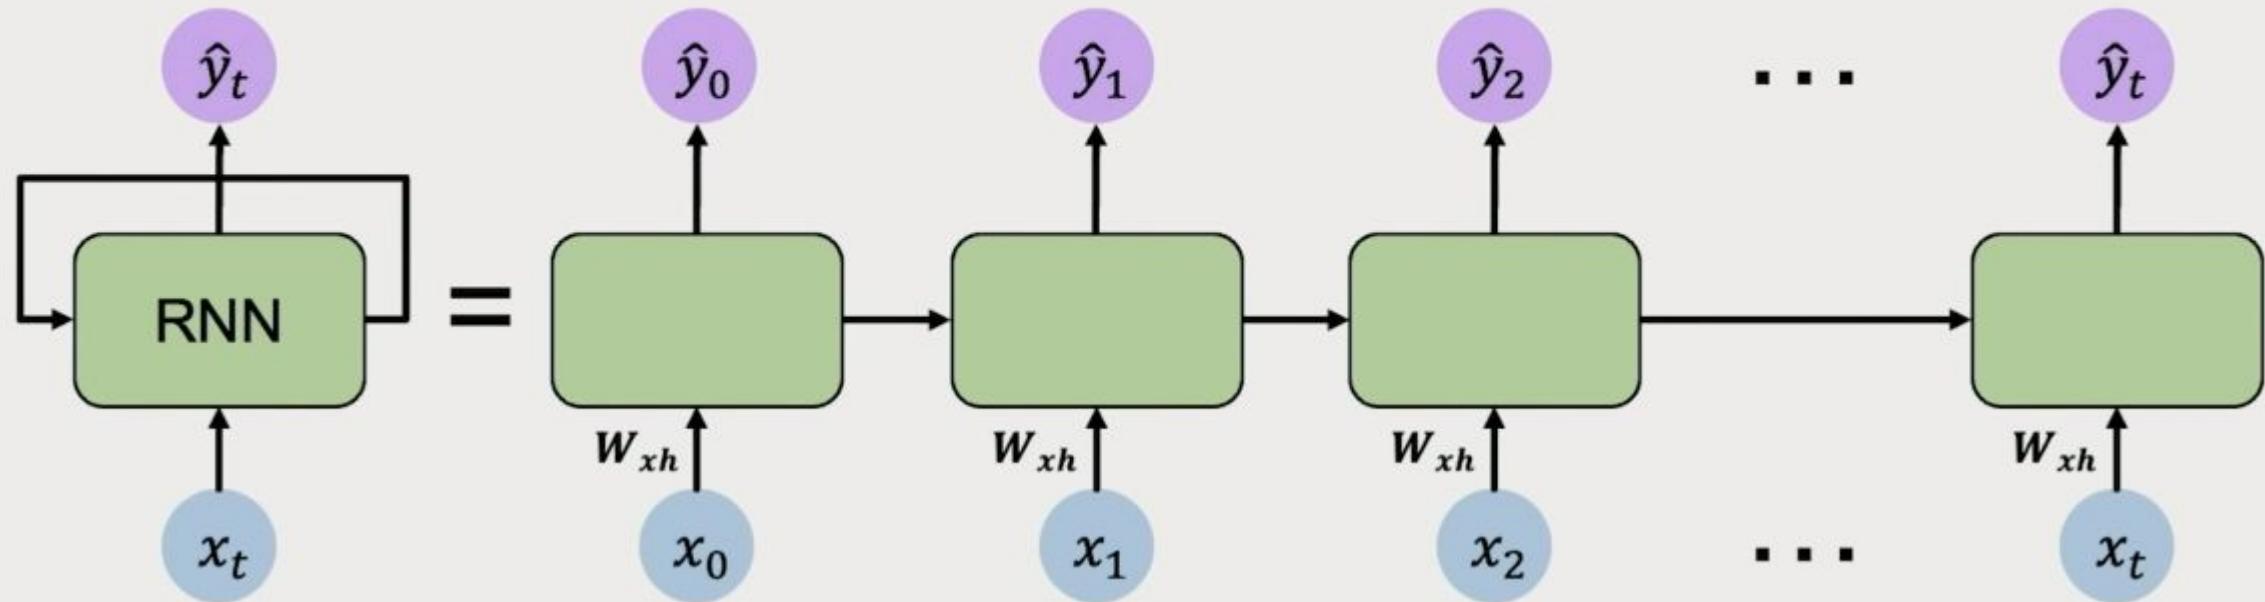
Input Vector

$$x_t$$

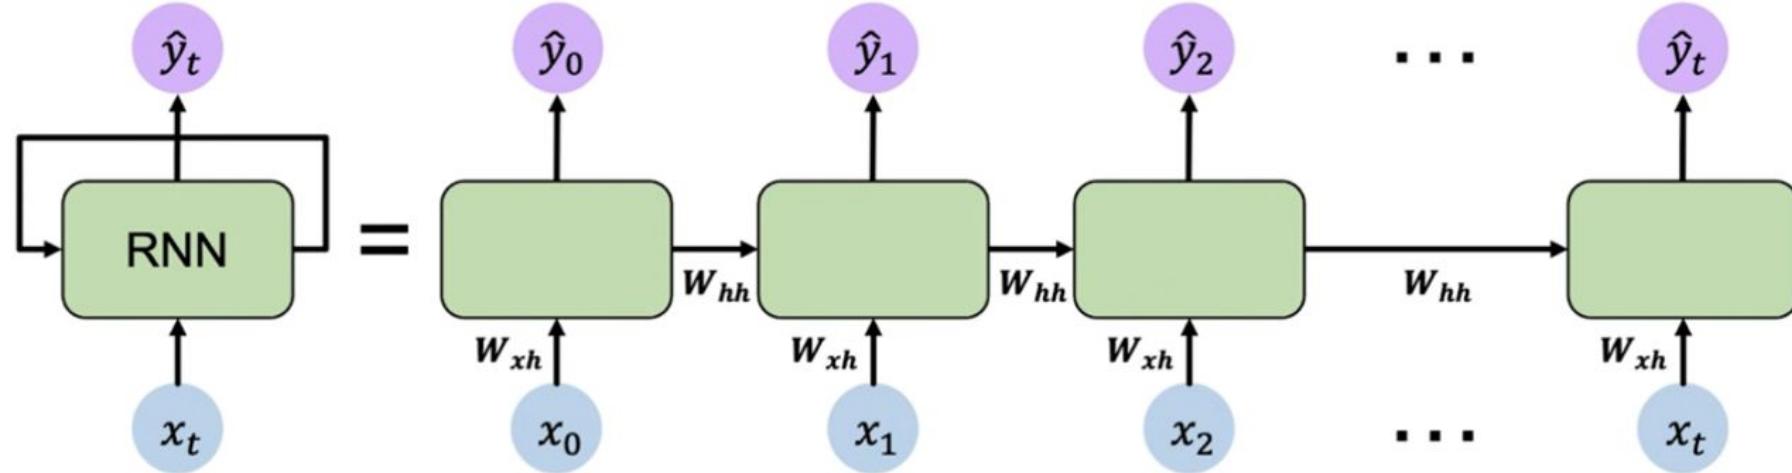
# RNN: Computational Graph Across Time



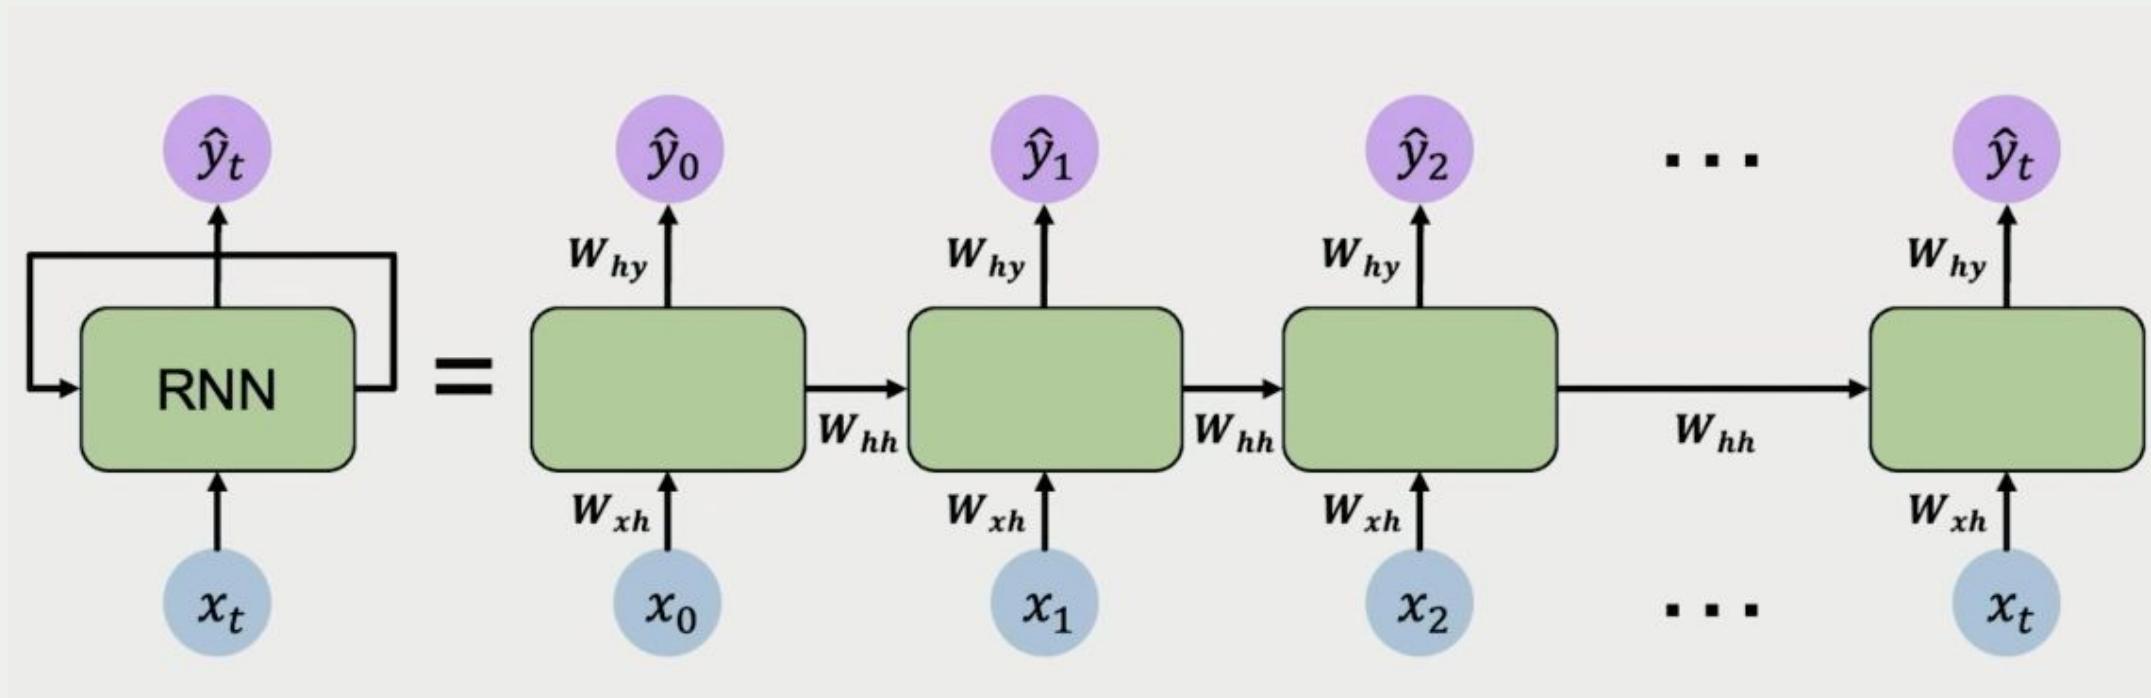
# RNN: Computational Graph Across Time



# RNN: Computational Graph Across Time



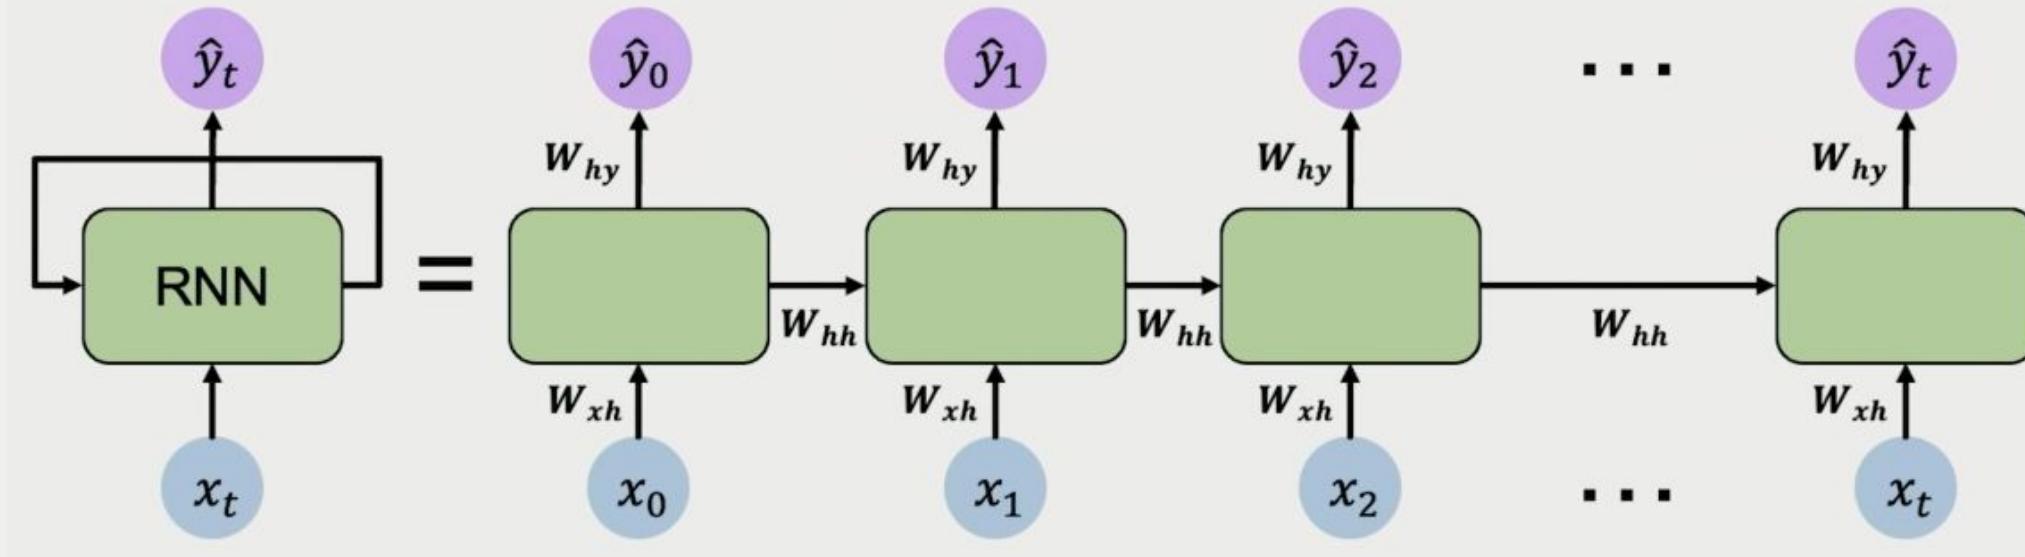
# RNN: Computational Graph Across Time



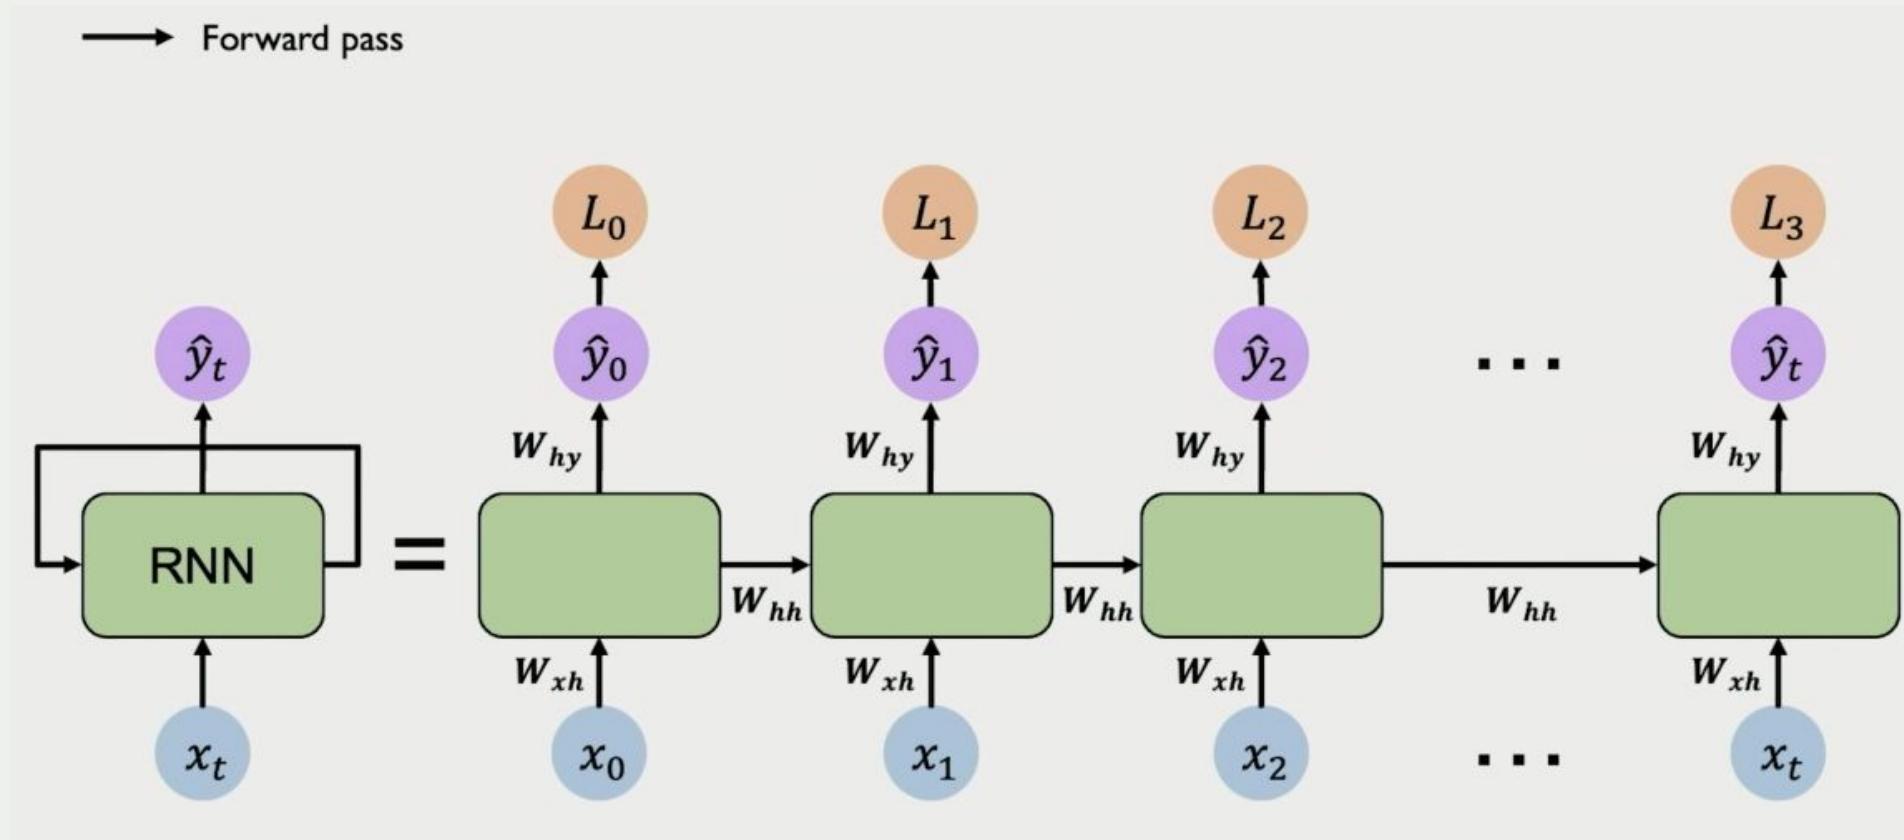
Reuse the same weight at every time step

# RNN: Computational Graph Across Time

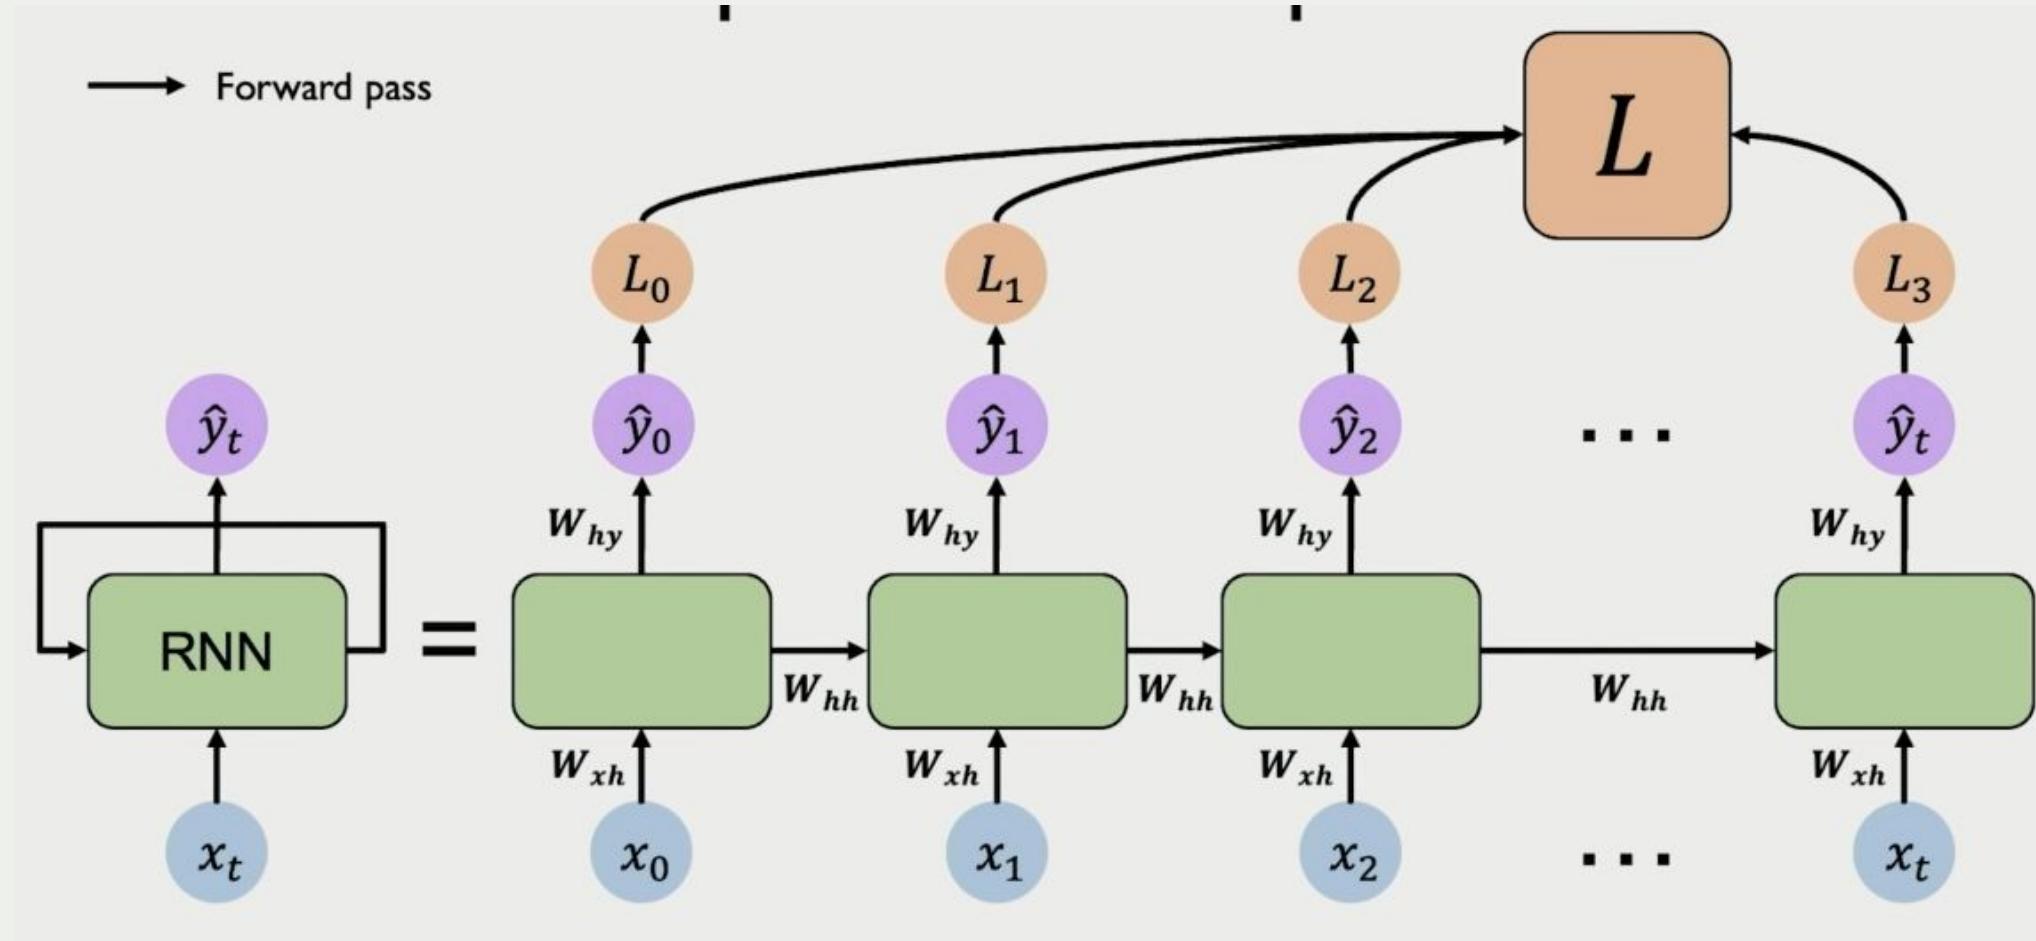
Re-use the **same weight matrices** at every time step



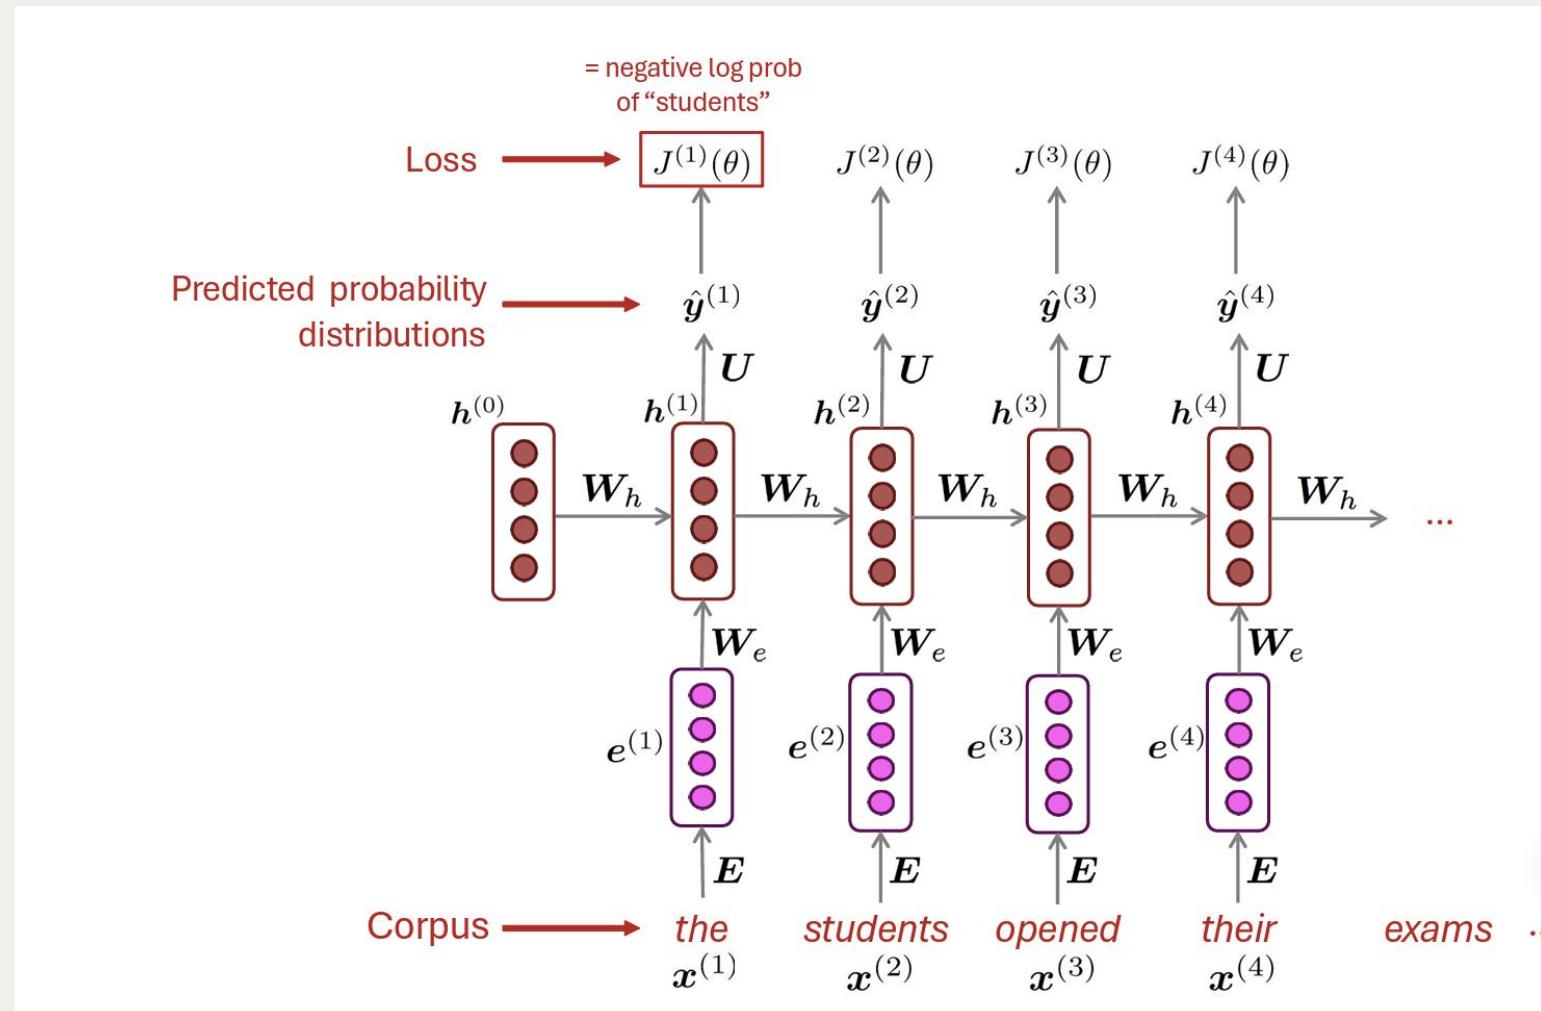
# RNN: Computational Graph Across Time



# RNN: Computational Graph Across Time



# Training an RNN Language Model

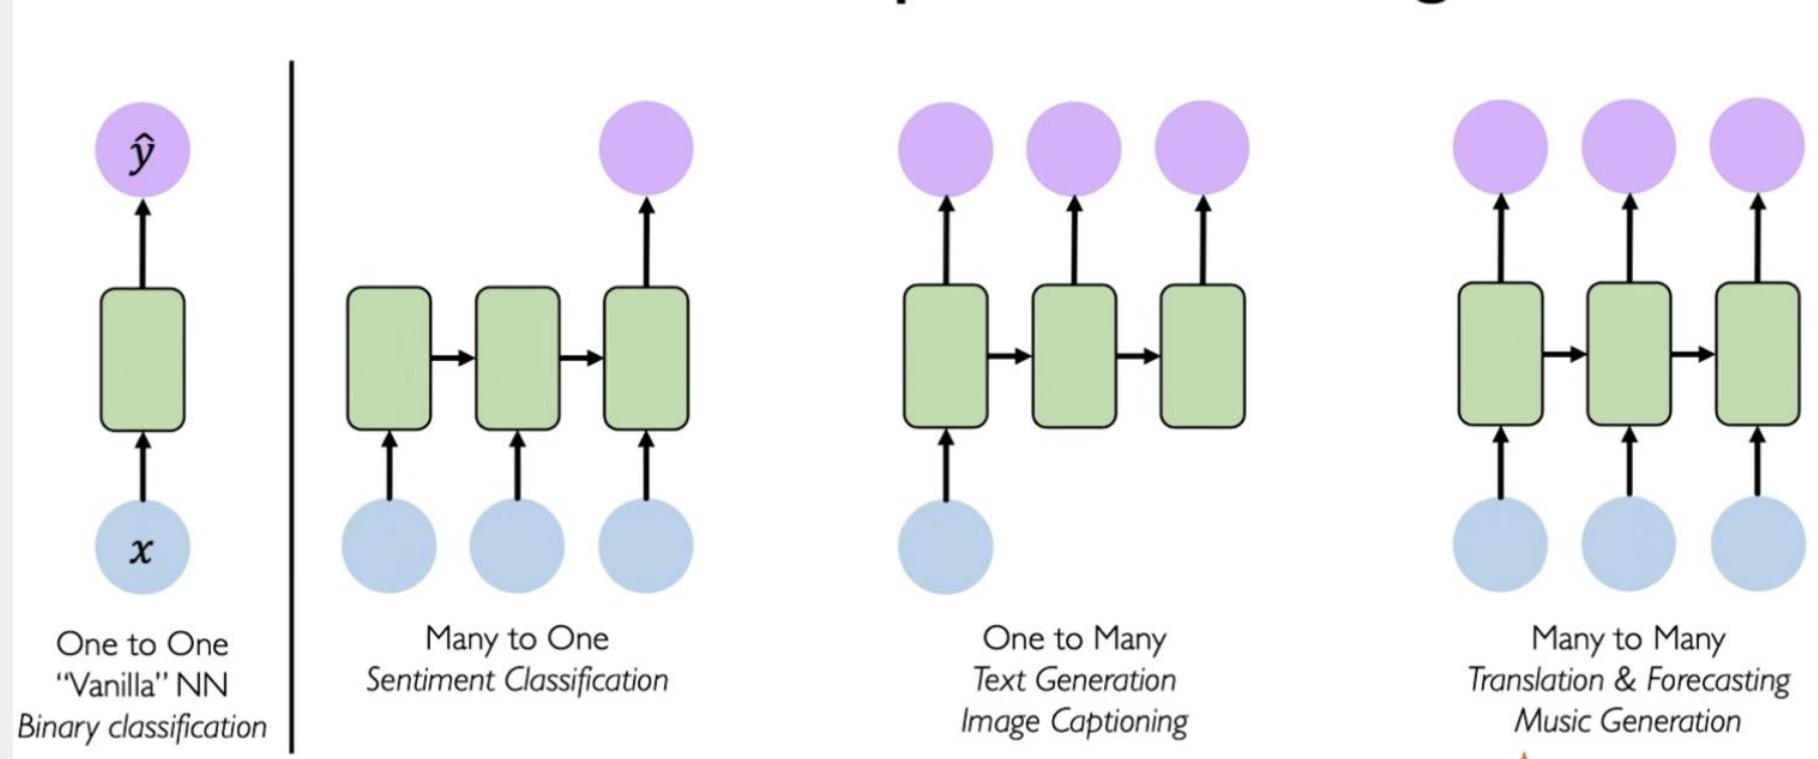


Chapter 6 Neural Networks: Speech and Language Processing. Daniel Jurafsky & James H. Martin.

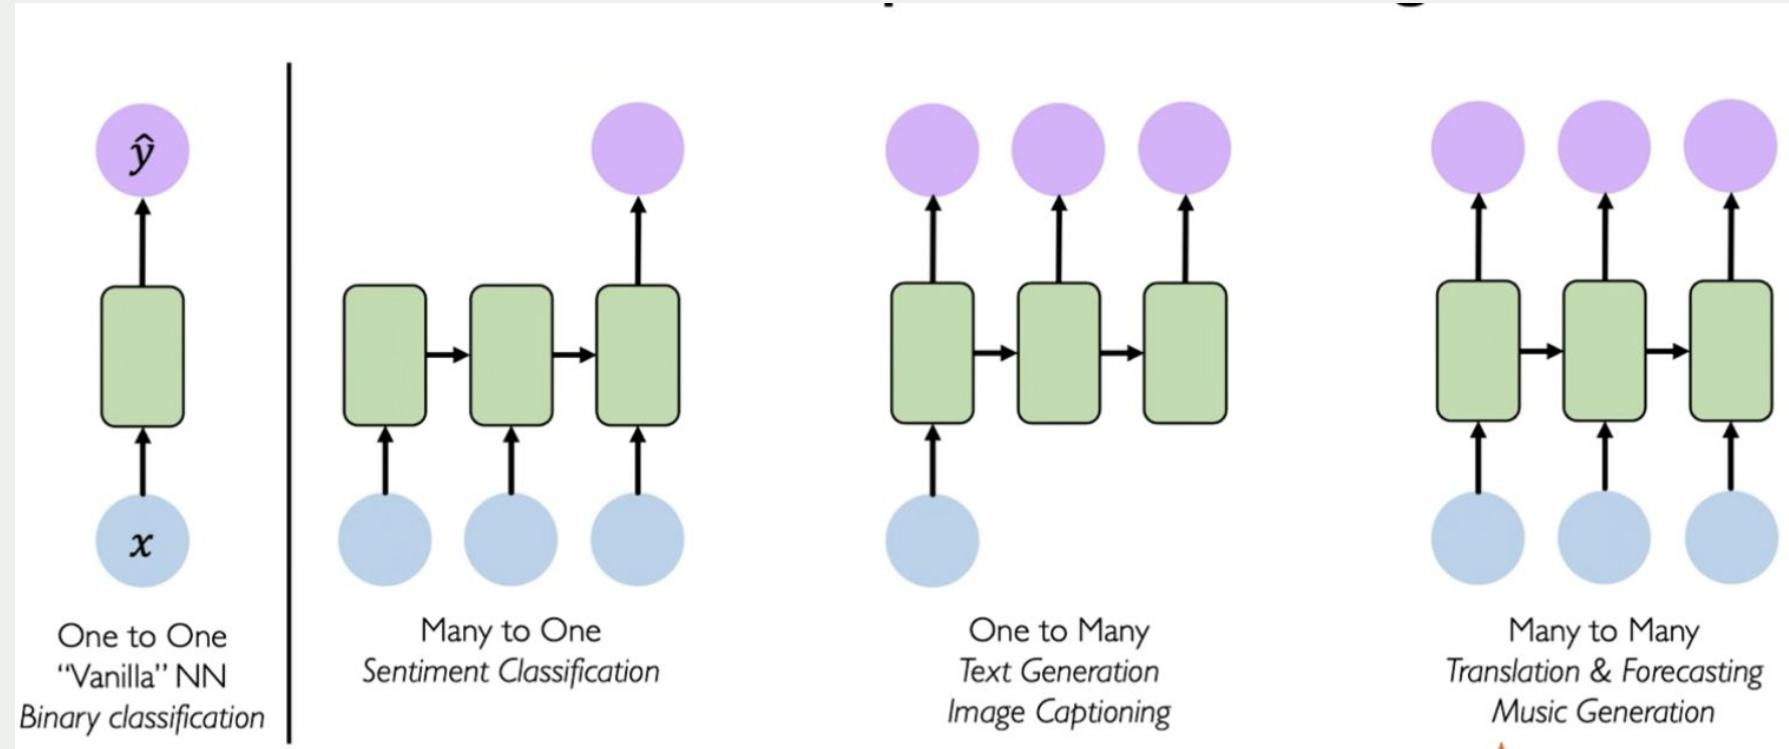
# Practical Session: RNN from Scratch

- CNN for text classification
- RNN for text classification

# RNNs for Sequence Modelling



# RNNs for Sequence Modelling

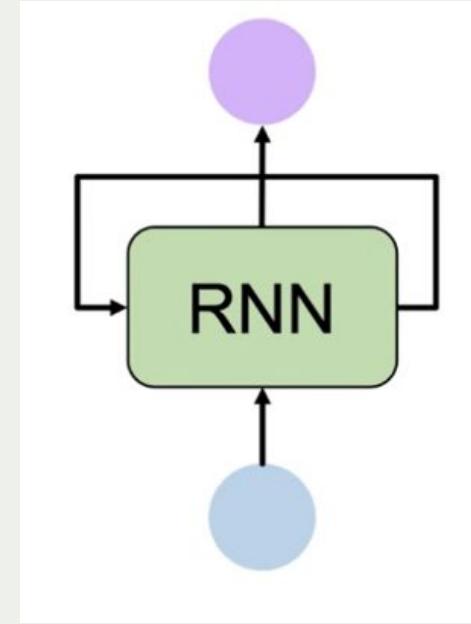


... and many other architectures and applications

# Sequence Modelling: Design Criteria

To model sequences, we need to:

- Handle variable-length sequences
- Track long-term dependencies
- Maintain information about order
- Share parameters across the sequence



**Recurrent Neural Networks (RNNs)** meet these sequence modeling design criteria.

Predict the next word

# A Sequence Modelling Problem

# A Sequence Modelling Problem

“This morning I took my cat for a walk.”

given these words

predict the  
next word

## Representing Language to a Neural Network

# A Sequence Modelling Problem

“This morning I took my cat for a walk.”

given these words

predict the  
next word

## Representing Language to a Neural Network



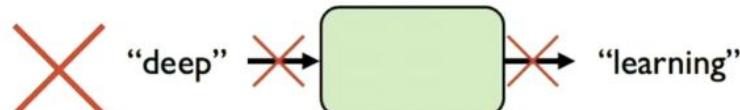
# A Sequence Modelling Problem

“This morning I took my cat for a walk.”

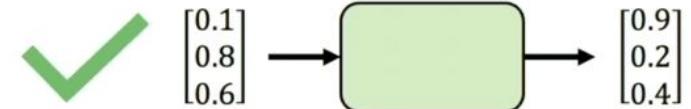
given these words

predict the  
next word

## Representing Language to a Neural Network

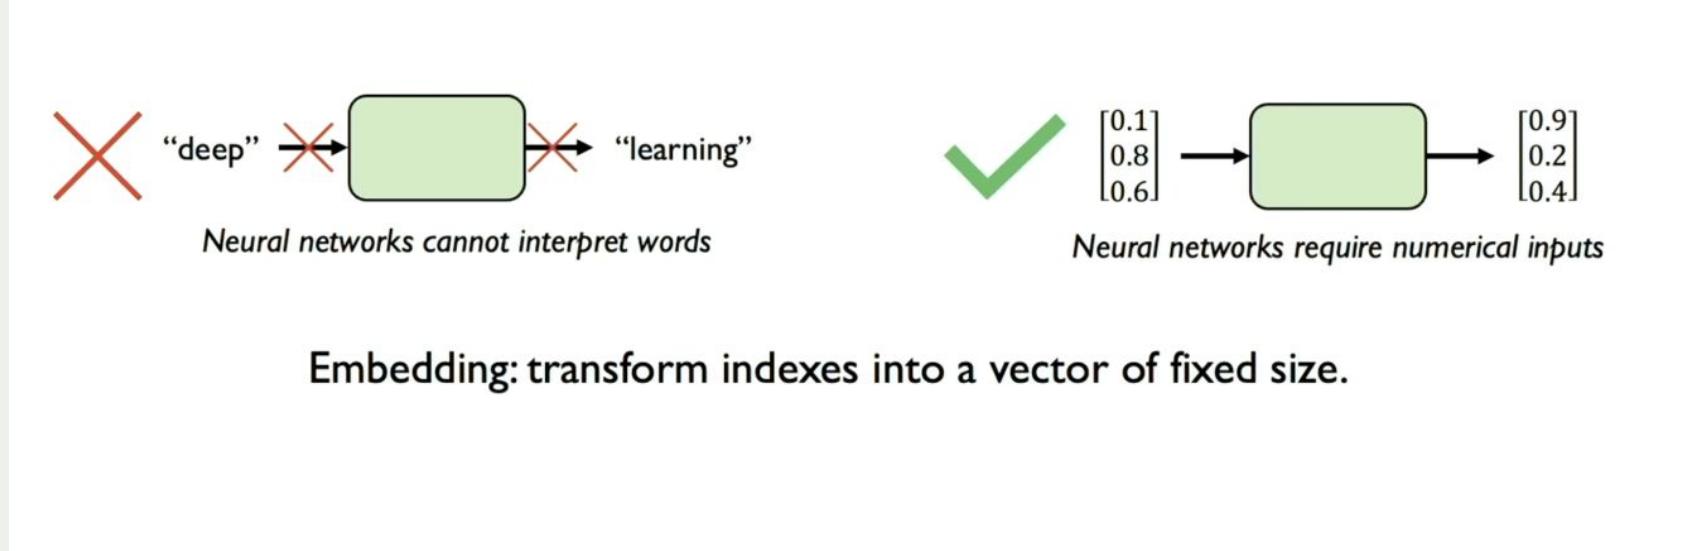


*Neural networks cannot interpret words*

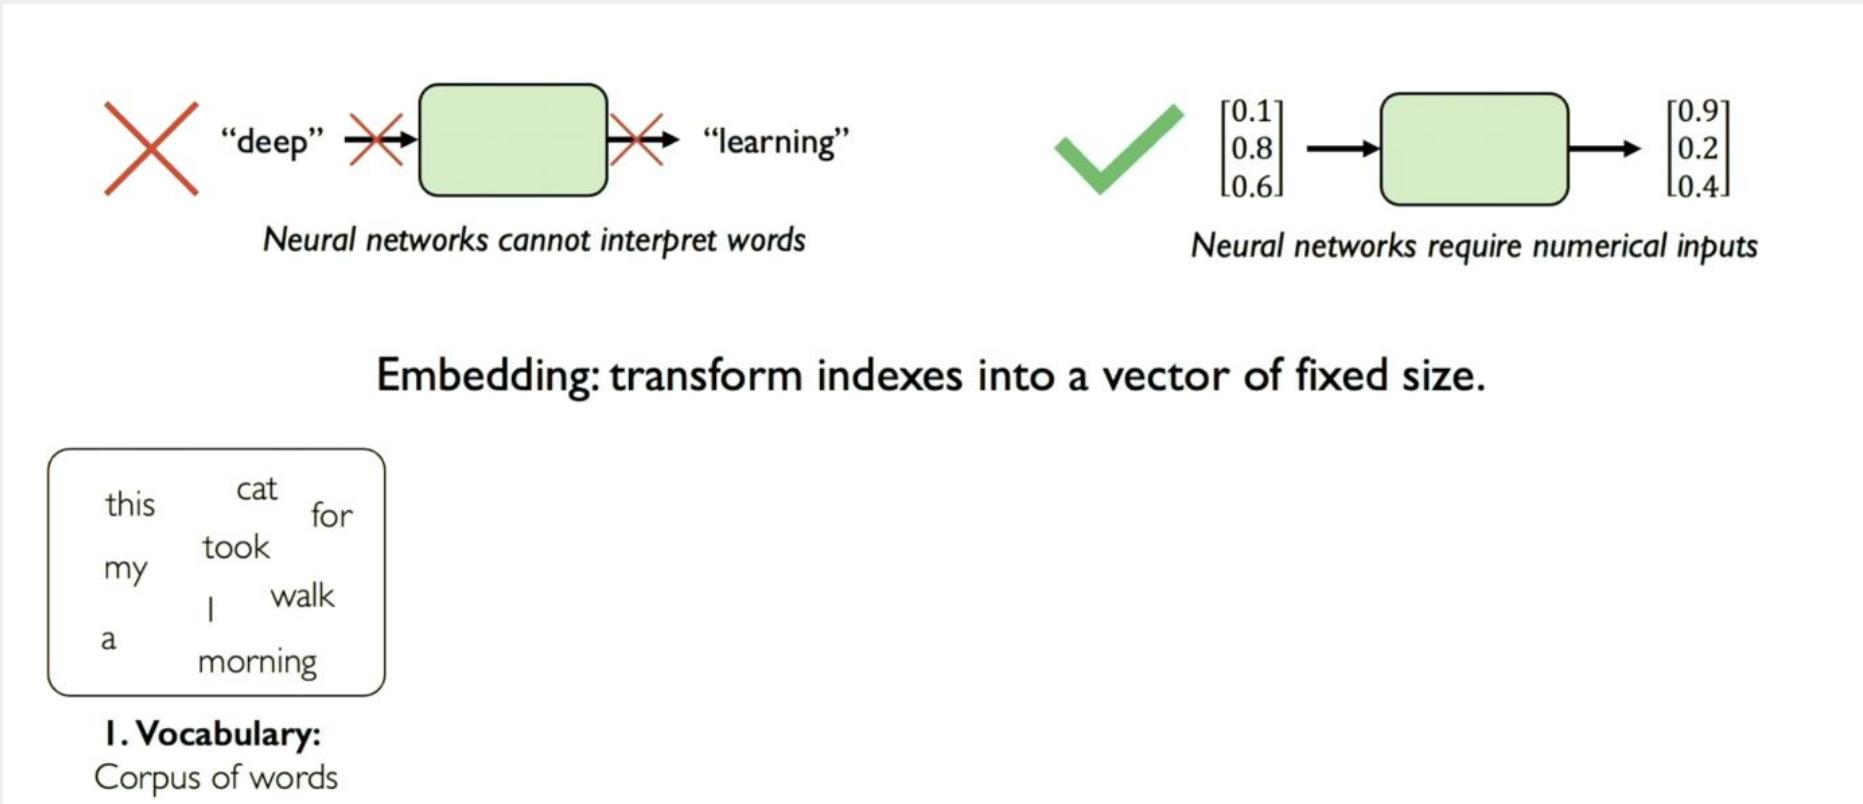


*Neural networks require numerical inputs*

# Encoding Language for a Neural Network



# Encoding Language for a Neural Network



# Encoding Language for a Neural Network



“deep” “learning”

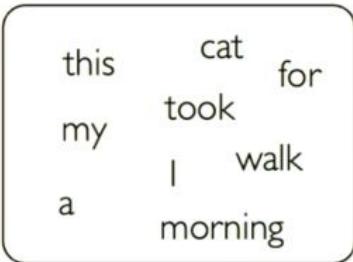
*Neural networks cannot interpret words*



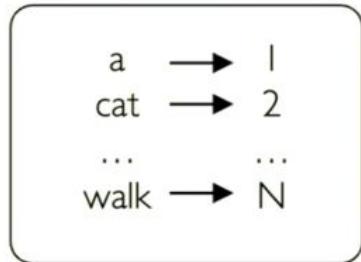
$\begin{bmatrix} 0.1 \\ 0.8 \\ 0.6 \end{bmatrix}$   $\begin{bmatrix} 0.9 \\ 0.2 \\ 0.4 \end{bmatrix}$

*Neural networks require numerical inputs*

**Embedding:** transform indexes into a vector of fixed size.



**1. Vocabulary:**  
Corpus of words

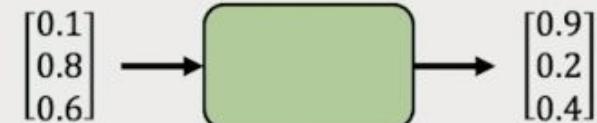


**2. Indexing:**  
Word to index

# Encoding Language for a Neural Network



*Neural networks cannot interpret words*



*Neural networks require numerical inputs*

**Embedding:** transform indexes into a vector of fixed size.

this            cat            for  
my            took             
                I              walk  
a                I              morning

**1. Vocabulary:**  
Corpus of words

a → 1  
cat → 2  
...  
walk → N

**2. Indexing:**  
Word to index

**One-hot embedding**  
“cat” = [ 0, 1, 0, 0, 0, 0 ]  
 $i$ -th index

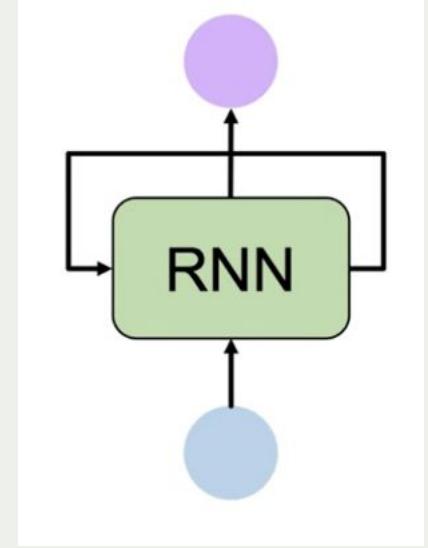
**Learned embedding**

**3. Embedding:**  
Index to fixed-sized vector

# Sequence Modelling: Design Criteria

To model sequences, we need to:

- **Handle variable-length sequences**
- Track **long-term** dependencies
- Maintain information about **order**
- **Share parameters** across the sequence



**Recurrent Neural Networks (RNNs)** meet these sequence modelling design criteria.

# Handle variable-length sequences

The food was great

vs.

We visited a restaurant for lunch

vs.

We were hungry but cleaned the house before eating

# Track long-term dependencies

"France is where I grew up, but I now live in Boston. I speak fluent \_\_\_\_."

We need information from **the distant past** to accurately predict the correct word.

# Capture Differences in Sequence Order



The food was good, not bad at all.

vs.

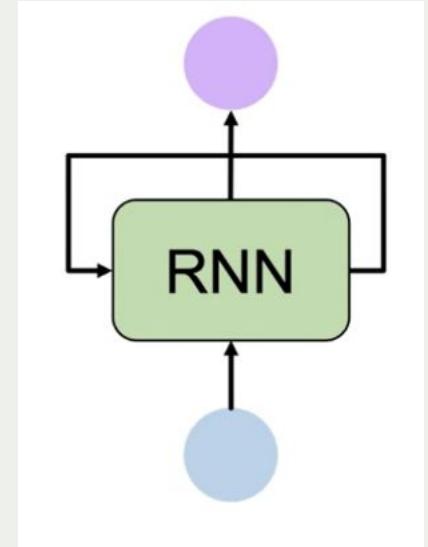
The food was bad, not good at all.



# Sequence Modelling: Design Criteria

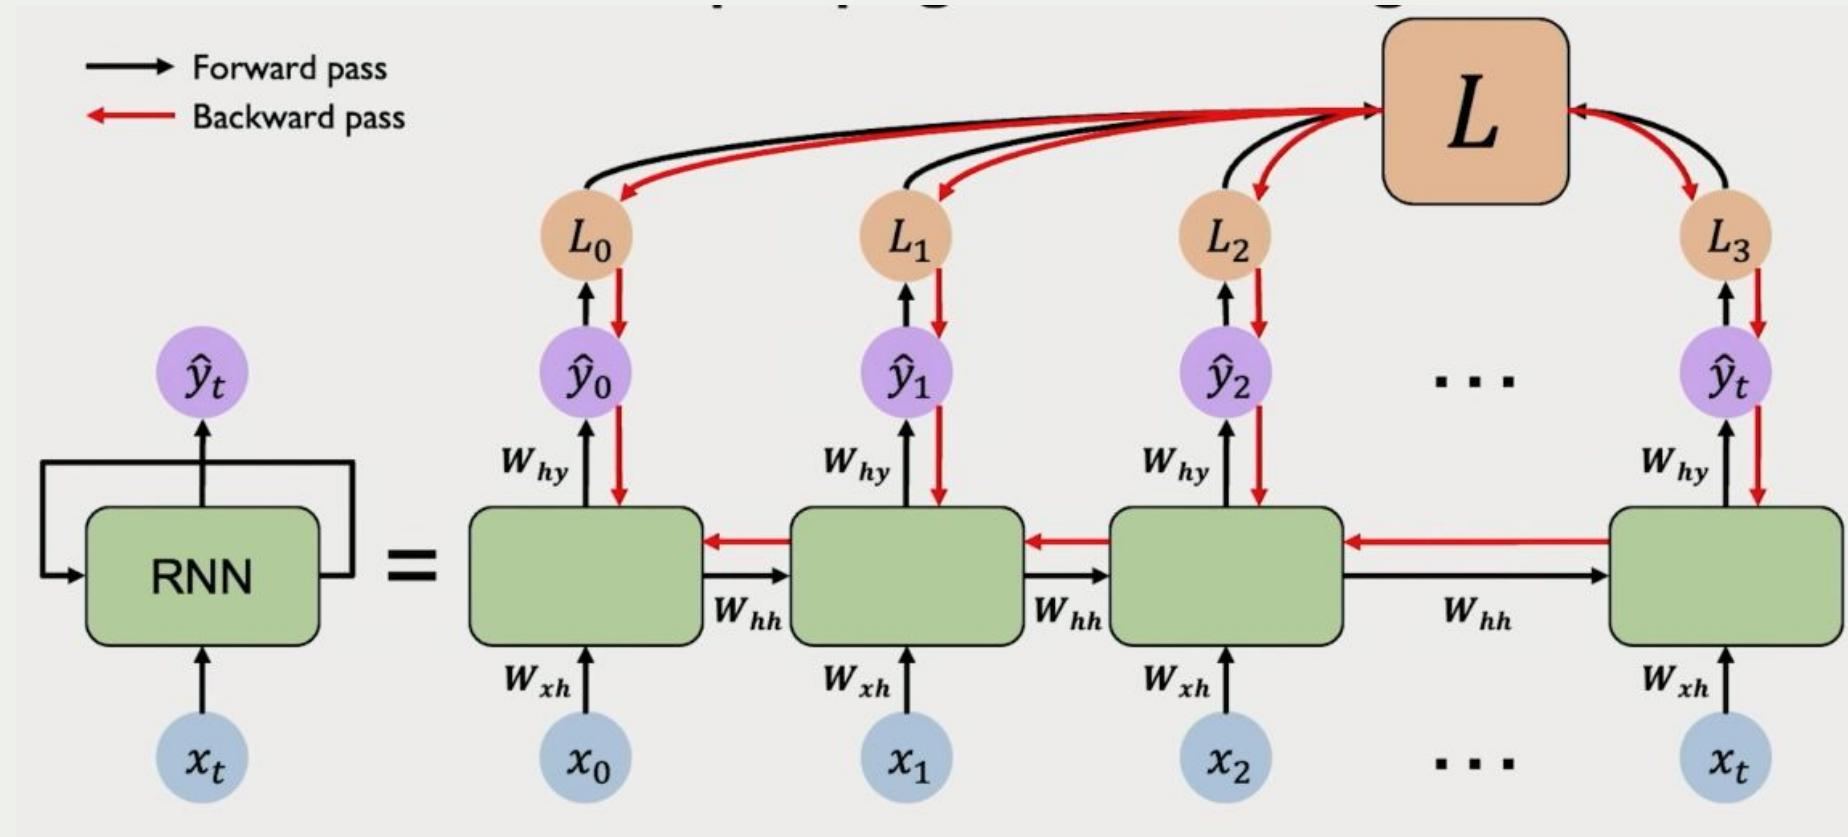
To model sequences, we need to:

- Handle variable-length sequences
- Track long-term dependencies
- Maintain information about order
- Share parameters across the sequence



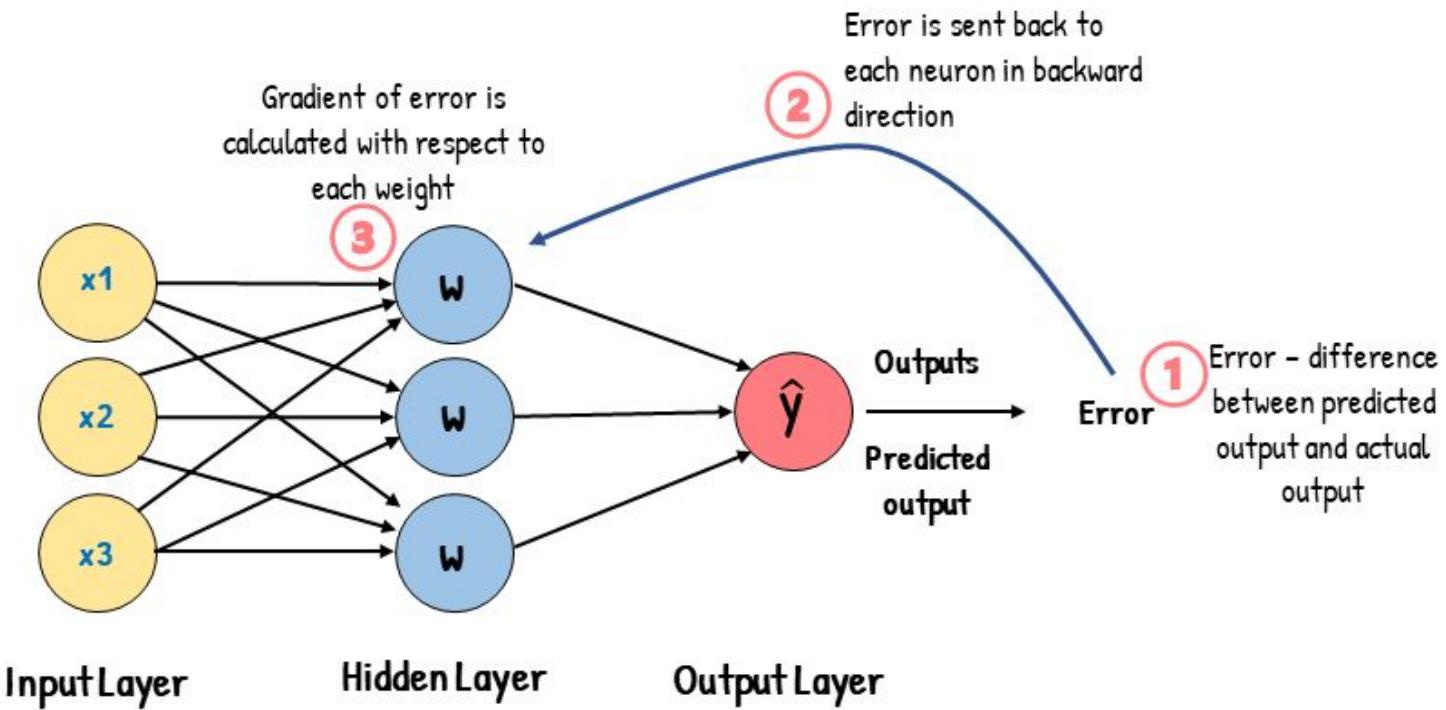
**Recurrent Neural Networks (RNNs)** meet these sequence modelling design criteria.

# Recall: Backpropagation in Feed Forward



# Recall: Backpropagation in Feed Forward

## Backpropagation



# Problems with RNNs

## Problem

- Vanishing Gradient
- Exploding Gradient
- Difficulty with Long-Term Memory
- Slow Training
- Limited Parallelization

# Q and A